



Reserve Bank of India

FREE-AI Committee Report

Framework for Responsible and Ethical Enablement
of Artificial Intelligence



7 Sutras

August 2025

Letter of Transmittal

Shri Sanjay Malhotra
Governor
Reserve Bank of India
Mumbai - 400 001

August 13, 2025

Dear Sir,

Report of the Committee - Framework for Responsible and Ethical Enablement of Artificial Intelligence in the Financial Sector

We are pleased to submit the report of the Committee constituted to develop a Framework for Responsible and Ethical Enablement of Artificial Intelligence (FREE-AI) in the financial sector. The Committee has approached its work with an open mind and has taken a holistic view of the opportunities and risks. The framework is anchored in seven sutras which serve as the foundational principles and is operationalised through twenty-six targeted recommendations under six strategic pillars. Together, they offer a forward-looking blueprint for all stakeholders.

We thank you for entrusting us with this responsibility and believe that the FREE-AI Framework will help harness the potential of AI in a way that preserves and promotes trust in the financial sector.

Yours sincerely,



(Dr. Pushpak Bhattacharyya)
Chairperson



(Ms. Debjani Ghosh)
Member



(Dr. Balaraman Ravindran)
Member



(Shri Abhishek Singh)
Member



(Shri Rahul Matthan)
Member



(Shri Anjani Rathor)
Member



(Shri Sree Hari Nagaralu)
Member



(Shri Suwendu Pati)
Member Secretary

Table of Contents

Acknowledgements	i
List of Abbreviations	ii
Executive Summary	iv
Chapter 1 – Introduction and Background	1
1.1 Evolution of Artificial Intelligence and Machine Learning	1
1.2 AI and ML in Financial Services	3
1.3 Constitution of the Committee	3
1.4 Terms of Reference	5
1.5 Methodology	5
1.6 Structure of the Report.....	6
Chapter 2 – AI in Finance: Opportunities and Challenges	7
2.1 Benefits and Opportunities	7
2.2 Emerging Risks and Sectoral Challenges	10
Chapter 3 – AI Policy Landscape and Insights from the Ecosystem	16
3.1 Global Policy Developments and Approaches	16
3.2 India’s Policy Environment and Developments	23
3.3 Insights from Surveys and Stakeholder Engagements	25
Chapter 4 – Building a Responsible and Ethical AI Framework	34
4.1 Trust as the Cornerstone	34
4.2 Enablers and Considerations for Advancing Trustworthy AI	35
4.3 The Seven Sutras - Guiding Principles	36
4.4 Principles to Practice - Recommendations.....	39
4.5 Conclusion - Weaving It All Together.....	67
Summary of Sutras and Recommendations	68
Annexure I – Interactions with Stakeholders by the Committee	75
Annexure II – Interactions with Stakeholders by the Secretariat	77
Annexure III – IndiaAI Mission: Strategy and Status	78
Annexure IV – AI Specific Enhancements in RBI Master Directions	80
Annexure V – Suggested Outline of Board Policy on AI	82
Annexure VI – AI Incident Reporting Form (Indicative Sample)	84
References	85
Glossary of Key Terms	87

Acknowledgements

The Committee is grateful to the Shri Sanjay Malhotra, Governor, Reserve Bank of India, for the opportunity to contribute to this important area at a crucial juncture in the evolution of technology in the financial sector. The Committee would like to express gratitude to Shri T. Rabi Sankar, Deputy Governor, RBI for his vision, insights, and valuable perspectives, that enriched the report. The Committee is also thankful to Shri P. Vasudevan, Executive Director, RBI for his guidance and support.

As part of the deliberations, the Committee engaged with a wide range of stakeholders to gain diverse perspectives on the adoption, opportunities, and challenges of artificial intelligence in the financial sector. The inputs were instrumental in developing a well-rounded understanding of the evolving AI ecosystem in India. The Committee is thankful for the interactions and acknowledges the contributions of all stakeholders who shared their time and expertise. A detailed list is provided in Annexure I.

The Committee would like to convey its appreciation to the Secretariat team of FinTech Department, comprising Shri Muralidhar Manchala, Shri Ankur Singh, Shri Praveen John Philip, Shri Padarabinda Tripathy, Shri Manan Nagori, Shri Ritam Gangopadhyay, for their excellent support in facilitating the Committee meetings and stakeholder interactions, conducting background research and survey, as well as assisting in the drafting of this report.

List of Abbreviations

A2A	Agent-to-Agent
AI	Artificial Intelligence
AIFI	All India Financial Institutions
AISI	Artificial Intelligence Safety Institute
BCP	Business Continuity Plan
DoS	Department of Supervision
DPDP Act	Digital Personal Data Protection Act
DPI	Digital Public Infrastructure
EDR	Endpoint Detection and Response
EmTech	Emerging Technology
ESG	Environmental, Social, and Corporate governance
FCA	Financial Conduct Authority
FREE-AI	Framework for Responsible and Ethical Enablement of Artificial Intelligence
FSB	Financial Stability Board
FSR	Financial Sector Regulator
FTD	FinTech Department
GenAI	Generative Artificial Intelligence
GPU	Graphics Processing Unit
IBA	Indian Banks' Association
IRDAI	Insurance Regulatory and Development Authority of India
ISO/IEC	International Organization for Standardization and International Electrotechnical Commission
KYC	Know Your Customer
LLM	Large Language Model
MCP	Model Context Protocol
MeitY	Ministry of Electronics and Information Technology
ML	Machine Learning
MRM	Model Risk Management
NABARD	National Bank for Agriculture and Rural Development

NBFC	Non-Banking Financial Company
NDSAP	National Data Sharing and Accessibility Policy
NLP	Natural Language Processing
OECD	Organisation for Economic Co-operation and Development
PET	Privacy-Enhancing Technologies
PIDF	Payment Infrastructure Development Fund
RAG	Retrieval Augmented Generation
RE	Regulated Entity
SCB	Scheduled Commercial Banks
SEBI	Securities and Exchange Board of India
SIEM	Security Information and Event Management
SLA	Service Level Agreement
SME	Small and Medium Enterprises
SLM	Small Language Model
SRO	Self-Regulatory Organisation
TSP	Technology Service Provider
UPI	Unified Payments Interface

Executive Summary

Artificial Intelligence (AI) is the transformative general-purpose technology of the modern age. Over the years, the simple rule-based models have evolved into complex systems capable of operating with limited human intervention. More recently, it has started to reshape how we work, how businesses operate and engage with their customers. In the process, it has forced us to question some of our most fundamental assumptions about human creativity, intelligence and autonomy.

For an emerging economy like India, AI presents new ways to address developmental challenges. Multi-modal, multi-lingual AI can enable the delivery of financial services to millions who have been excluded. **When used right, AI offers tremendous benefits. If used without guardrails, it can exacerbate the existing risks and introduce new forms of harm.**

The challenge with regulating AI is in striking the right balance, making sure that society stands to gain from what this technology has to offer, while mitigating its risks. Jurisdictions have adopted different approaches to AI policy and regulation based on their national priorities and institutional readiness.

In the financial sector, AI has the potential to unlock new forms of customer engagement, enable alternate approaches to credit assessment, risk monitoring, fraud detection, and offer new supervisory tools. At the same time, increased adoption of AI could lead to new risks like bias and lack of explainability, as well as amplifying existing challenges to data protection, cybersecurity, among others.

In order to encourage the responsible and ethical adoption of AI in the financial sector, the FREE-AI Committee was constituted by the Reserve Bank of India. The RBI conducted two surveys to understand current AI adoption and challenges in the financial sector. The Committee referenced these surveys and, in addition, undertook extensive stakeholder consultations to gain further insights.

After extensive deliberations, the Committee formulated **7 Sutras** that represent the core principles to guide AI adoption in the financial sector. These are:

- (i) **Trust is the Foundation**
- (ii) **People First**
- (iii) **Innovation over Restraint**

- (iv) Fairness and Equity
- (v) Accountability
- (vi) Understandable by Design
- (vii) Safety, Resilience and Sustainability

Using the *Sutras* as guidance, the Committee recommends an approach that fosters innovation and mitigates risks, treating these two seemingly competing objectives as complementary forces that must be pursued in tandem. This is achieved through a unified vision spread across **6 strategic Pillars** that address the dimensions of innovation enablement as well as risk mitigation. Under innovation enablement, the focus is on **Infrastructure, Policy** and **Capacity** and for risk mitigation, the focus is on **Governance, Protection** and **Assurance**. Under these six pillars, the report outlines **26 Recommendations** for AI adoption in the financial sector.

To foster innovation, it recommends:

- the establishment of shared infrastructure to democratise access to data and compute; the creation of an AI Innovation Sandbox
- the development of indigenous financial sector-specific AI models
- the formulation of an AI policy to provide necessary regulatory guidance
- institutional capacity building at all levels, including the board and the workforce of REs and other stakeholders,
- the sharing of best practices and learnings across the financial sector
- a more tolerant approach to compliance for low-risk AI solutions to facilitate inclusion and other priorities

To mitigate AI risks, it recommends:

- the formulation of a board-approved AI policy by REs
- the expansion of product approval processes, consumer protection frameworks and audits to include AI related aspects
- the augmentation of cybersecurity practices and incident reporting frameworks
- the establishment of robust governance frameworks across the AI lifecycle
- making consumers aware when they are dealing with AI

This is the **FREE-AI vision**: a financial ecosystem where the encouragement of innovation is in harmony with the mitigation of risk.

FREE AI Framework

INNOVATION ENABLEMENT

Infrastructure

1. Financial Sector Data Infrastructure
2. AI Innovation Sandbox
3. Incentives and Funding Support
4. Indigenous Financial Sector-Specific AI Models
5. Integrating AI with DPI

Policy

6. Adaptive and Enabling Policies
7. Enabling AI-Based Affirmative Action
8. AI Liability Framework
9. AI Institutional Framework

Capacity

10. Capacity Building within REs
11. Capacity Building for Regulators and Supervisors
12. Framework for Sharing Best Practices
13. Recognise and Reward Responsible AI Innovation

RISK MITIGATION

Governance

14. Board Approved AI Policy
15. Data Lifecycle Governance
16. AI System Governance Framework
17. Product Approval Process

Protection

18. Consumer Protection
19. Cybersecurity Measures
20. Red-Teaming
21. Business Continuity Plan for AI Systems
22. AI Incident Reporting and Sectoral Risk Intelligence Framework

Assurance

23. AI Inventory within REs and Sector-wide Repository
24. AI Audit Framework
25. Disclosures by REs
26. AI Toolkit

7 Sutras

TRUST IS THE FOUNDATION

Trust is non-negotiable and should remain uncompromised.



PEOPLE FIRST

AI should augment human decision-making but defer to human judgment and citizen interest.

INNOVATION OVER RESTRAINT

Foster responsible innovation with purpose.



FAIRNESS AND EQUITY

AI outcomes should be fair and non-discriminatory

ACCOUNTABILITY

Accountability rests with the entities deploying AI



UNDERSTANDABLE BY DESIGN

Ensure explainability for trust

SAFETY, RESILIENCE, AND SUSTAINABILITY

AI systems should be secure, resilient and energy efficient



Chapter 1 – Introduction and Background

“It is not the strongest of the species that survive, but the most adaptable to change.”

- Charles Darwin

Artificial Intelligence (AI) has seen significant growth in recent years, drawing attention from industry, innovators, policy makers and consumers alike. Whether it is seeking answers, creating avatars, or personalised e-commerce, AI is increasingly getting embedded in day-to-day activities. Given the recent surge in interest, it is easy to view AI as a relatively new phenomenon. However, the roots of AI actually date back several decades.

1.1 Evolution of Artificial Intelligence and Machine Learning

1.1.1 Early Foundations and Milestones: In his seminal 1950 paper *Computing Machinery and Intelligence*, renowned mathematician Alan Turing first posed the fundamental question, “*Can machines think?*” and then introduced the Imitation Game (now known as the Turing Test) as a way to gauge machine intelligence. However, the term “Artificial Intelligence” was coined in 1956 by John McCarthy during the Dartmouth Summer Research Project on Artificial Intelligence, a seminal event which set the stage for decades of exploration.

1.1.2 Early research in the 1960s and 1970s focused on symbolic AI and logic-based programs (the era of “Good Old-Fashioned AI” (GOFAI)) that could prove mathematical theorems and solve puzzles. These periods of over-optimism were followed by “AI winters” when funding and interest waned, however, foundational work continued. By the 1980s, expert systems, i.e., rule-based programs encoding human expert knowledge, became popular. Yet, these systems were hard to maintain and required manual knowledge engineering.

1.1.3. Emergence of Machine Learning: Machine Learning (ML) enabled algorithms to learn autonomously from data without explicit programming. This shift in the 1990s was due to significant improvements in computing power, data storage, and connectivity. ML techniques like neural networks, decision trees, and support vector machines began outperforming rule-based systems in tasks like image classification and language translation. World Chess Champion Garry Kasparov’s 3½ - 2½ defeat

to IBM's Deep Blue in a six-game rematch in 1997 demonstrated the ability of machines to outperform humans in domains considered to require strategic reasoning. This inspired early exploration in financial applications as well.

1.1.4 As a financial sector application, HNC Software's Falcon system was screening two-thirds of all credit card transactions worldwide by the 1990s. ML application grew in the 2000s, and in finance, early ML models were deployed for specific, well-defined tasks: for instance, using neural networks, Banks also adopted ML for credit scoring beyond traditional logistic regression, using larger datasets to enhance prediction accuracy.

1.1.5 The Deep Learning Revolution and Generative AI: The 2010s saw further breakthroughs with the rise of deep learning, a subset of ML that involved multi-layered neural networks. A major milestone during this period was the release of the 2017 paper "*Attention is All You Need*" by researchers at Google, which introduced the Transformer architecture that laid the foundation for large language models (LLMs). The power of deep learning's ability to carry out complex pattern recognition was validated by landmark achievements such as computers surpassing human accuracy in image recognition in 2012 and when Google DeepMind's "AlphaGo" defeated Go champion Lee Sedol in 2016. Soon after, voice assistants became commonplace, and self-driving cars took to the roads. AI was no longer confined to labs; it began to surface in everyday products and services.

1.1.6 In late 2022, Generative AI tools brought the power of advanced AI directly to the public. ChatGPT reached 100 million users in just two months after launch¹, highlighting the unprecedented pace of adoption. Techniques such as retrieval-augmented generation (RAG), mixture-of-experts (MoE) architectures are further enhancing capabilities. From generating images to creating complex reports using a suite of agents, AI has moved beyond just being a niche technology to gradually reshaping the way we work.

1.1.7 Unprecedented Progress: As per the AI Index report 2025 by Stanford, AI systems now outperform humans in nearly all tested domains. Complex reasoning is the last major frontier, but even here, the gap is narrowing quickly. Open-source AI

¹ <https://www.reuters.com/technology/chatgpt-sets-record-fastest-growing-user-base-analyst-note-2023-02-01/>

models are rapidly catching up to closed models, narrowing the gap from 8% to just 1.7%. Smaller models are also showing significant gains in efficiency and capability. The year 2024 marked a shift in national strategy with record public investments: India (\$1.25 billion), France (\$117 billion), Canada (\$2.40 billion), China (\$47.50 billion), and Saudi Arabia (\$100 billion)².

1.2 AI and ML in Financial Services

1.2.1 The role of AI in financial services has significantly increased over the last decade. As machine learning has matured, banks and insurers have expanded use cases from rule-based systems to real-time fraud detection, anomaly detection in claims processing, and market forecasting. The 2010s saw the rise of big data and deep learning, enabling institutions to leverage alternative data sources (e.g., social media, geolocation) and deploy NLP-powered chatbots like Bank of America's "Erica." Today, Gen-AI is being used in advanced chatbots, automated report generation, and the creation of synthetic data sets for safer model training. It is estimated that this could add \$200-340 billion annually to the global banking sector through productivity gains in compliance, risk management, and customer service³.

1.2.2 In the Indian context, AI has the potential to improve financial inclusion, expand opportunities for innovation and enhance efficiency in financial systems. Yet, these systems pose certain incremental risks and ethical dilemmas. As these systems are being increasingly integrated into high-stakes applications such as credit approvals, fraud detection, and compliance, there is a need to ensure that their application is responsible and ethical, that harm does not arise from their use, and that their outcomes do not undermine public trust.

1.3 Constitution of the Committee

1.3.1 In order to further responsible innovation in AI, while at the same time ensuring that consumer interests are protected, the Reserve Bank of India announced the establishment of a Committee to develop a framework for the responsible and ethical enablement of AI in the financial sector in its Statement on Developmental and

² Stanford: Artificial Intelligence Index Report 2025

³ <https://www.mckinsey.com/industries/financial-services/our-insights/scaling-gen-ai-in-banking-choosing-the-best-operating-model>

Regulatory Policies dated December 6, 2024⁴. Accordingly, the committee for developing the Framework for Responsible and Ethical Enablement of Artificial Intelligence in the Financial Sector (hereinafter referred to as the Committee or FREE-AI Committee) was constituted. The members of the committee are:

Sl. No.	Name	Member
i)	Dr. Pushpak Bhattacharyya, Professor, Department of Computer Science and Engineering, IIT Bombay	Chairperson
ii)	Ms. Debjani Ghosh, Distinguished Fellow, NITI Aayog; Independent Director, Reserve Bank Innovation Hub; and Ex-President, NASSCOM	Member
iii)	Dr. Balaraman Ravindran, Professor and Head, Wadhvani School of Data Science and AI, IIT Madras	Member
iv)	Shri Abhishek Singh, Additional Secretary, Ministry of Electronics and Information Technology, Government of India	Member
v)	Shri Rahul Matthan, Partner, Trilegal	Member
vi)	Shri Anjani Rathor, Group Head and Chief Digital Experience Officer, HDFC Bank Ltd.	Member
vii)	Shri Sree Hari Nagaralu, Head of Security AI Research, Microsoft India (R&D)	Member
viii)	Shri Suvendu Pati, CGM, FinTech Department, Reserve Bank of India	Member Secretary

⁴ https://www.rbi.org.in/Scripts/BS_PressReleaseDisplay.aspx?prid=59245

1.4 Terms of Reference

1.4.1 The terms of reference of the Committee are as under:

- i. To assess the current level of adoption of AI in financial services globally and in India.
- ii. To review regulatory and supervisory approaches on AI with a focus on the financial sector globally.
- iii. To identify potential risks associated with AI, if any, and recommend an evaluation, mitigation and monitoring framework and consequent compliance requirements for financial institutions, including banks, NBFCs, FinTechs, PSOs, etc.
- iv. To recommend a framework including governance aspects for responsible, ethical adoption of AI models/ applications in the Indian financial sector.
- v. Any other matter related to AI in the Indian financial sector.

1.5 Methodology

1.5.1 The Committee adopted a four-pronged approach.

i. Stakeholder Engagement: The Committee held extensive deliberations and adopted a consultative approach to get insights on the emerging developments, ongoing innovations, stakeholder needs, challenges and risks in the financial sector on account of the use of AI. Interactions were also conducted with stakeholders, including presentations from the RBI departments, consultants, and financial sector entities. Details of the interactions are provided at Annexure I and II.

ii. Survey and Interactions: Two targeted surveys were carried out, covering Scheduled Commercial Banks (SCBs), Non-Banking Financial Companies (NBFCs), All India Financial Institutions (AIFI) and FinTechs. Follow-up interactions were conducted with select Chief Digital Officers / Chief Technology Officers (CDOs/CTOs) to understand the extent to which AI had been adopted in the Indian financial services industry and any associated challenges.

iii. Review of global developments and literature: The Committee also examined the internationally published literature, global developments, extant regulatory frameworks/ approaches adopted in other jurisdictions and views of global standard-setting bodies (SSBs) and international organisations (IOs).

iv. Analysis of extant regulatory guidelines: Finally, the Committee analysed the extant regulatory framework applicable to the REs, such as those related to cybersecurity, data protection, consumer protection, and outsourcing, to the extent they capture the AI-specific risks and concerns.

1.5.2 In addition, based on the stakeholder engagement and survey feedback, the Committee acknowledged the need to place specific emphasis on fostering AI innovation and treated it as a critical reference point in defining its approach.

1.6 Structure of the Report

1.6.1 The remainder of the report is structured into three chapters. Chapter 2 examines the current state of AI adoption in the financial sector, highlighting the benefits and opportunities, and the evolving landscape of risks and challenges associated with AI deployment. Chapter 3 analyses the broader policy environment, covering key global approaches, domestic developments, and practical insights drawn from stakeholder interactions and survey responses across regulated entities and FinTechs. Finally, Chapter 4 presents the Committee's proposed Framework for Responsible and Ethical Enablement of Artificial Intelligence (FREE-AI). The terms used in this Report are explained in the Glossary at the end of this Report for contextual understanding.

Chapter 2 – AI in Finance: Opportunities and Challenges

“We can only see a short distance ahead, but we can see plenty there that needs to be done.” - Alan Turing, Computing Machinery and Intelligence, 1950

The financial services sector has witnessed the gradual integration of AI into core business functions such as risk management, fraud detection, and customer service. The recent AI evolution, while opening new frontiers of innovation, also gives rise to certain challenges about unintended outcomes and consequences. This chapter highlights the opportunities it offers and new risks that warrant more careful consideration.

2.1 Benefits and Opportunities

2.1.1 The adoption of AI in financial services has accelerated globally. According to a 2025 World Economic Forum white paper⁵ on AI in Financial Services, projected investments across banking, insurance, capital markets and payments business are expected to reach over ₹8 lakh crore (\$97 billion) by 2027. It is believed that AI will directly contribute to revenue growth in the coming years. The generative AI segment alone is forecast to cross ₹1.02 lakh crore (\$12 billion) by 2033, with a compound annual growth rate (CAGR) of 28–34%⁶. The OECD highlighted that AI is currently being developed or deployed by a broad range of financial institutions with major use cases such as customer relations, process automation and fraud detection⁷.

2.1.2 As AI continues to gain traction across financial services, it is beginning to unlock value by enhancing efficiency, accuracy and personalisation at scale. A key set of drivers underpinning this adoption includes the need to enhance customer experience, improve employee productivity, increase revenue, reduce operational costs, ensure regulatory compliance, and enable the development of new and innovative products. GenAI is poised to improve banking operations in India by up to 46%⁸. AI-driven analytics allow institutions to better understand customer behaviour, manage risk

⁵ https://reports.weforum.org/docs/WEF_Artificial_Intelligence_in_Financial_Services_2025.pdf

⁶ <https://www.statista.com/statistics/1449285/global-generative-ai-in-financial-services-market-size/>

⁷ https://www.oecd.org/en/publications/regulatory-approaches-to-artificial-intelligence-in-finance_f1498c02-en.html

⁸ Ernst & Young: How much productivity can GenAI unlock in India? The Aldea of India: 2025

proactively, and optimize operational costs. AI-powered alternate credit scoring models continue to expand credit access to the underserved population. AI chatbots can handle routine customer queries with 24x7 availability. AI-based early warning signals facilitate enhanced risk management. For instance, J.P.Morgan claims AI has significantly reduced fraud by improving payment validation screening, leading to a 15-20% reduction in account validation rejection rates and significant cost savings⁹. AI also improves operational efficiency through automating repetitive tasks such as data entry, document summarisation, and aiding human decisions.

2.1.3 AI for Financial Inclusion: In developing economies like India, where millions remain outside the ambit of formal finance, AI can help assess creditworthiness using non-traditional data sources such as utility payments, mobile usage patterns, GST filings, or e-commerce behaviour, thereby including “thin-file” or “new-to-credit” borrowers. AI-powered chatbots can offer context-aware financial guidance, grievance redressal, and behavioural nudges to low-income and rural populations. Voice-enabled banking in regional languages has the potential to allow illiterate or semi-literate individuals to access finance.

2.1.4 Leveraging AI in Digital Public Infrastructure: The 2023 recommendations of the G20 Task Force on DPI¹⁰ highlighted the need to integrate AI responsibly with DPI. India’s pioneering DPI ecosystem, including Aadhaar, UPI frameworks, offers a robust foundation for AI-driven enhanced service delivery, personalisation and real-time decision making. This convergence can pave the path for next-gen DPI where services are not only digital, but intelligent, inclusive and adaptive. Conversational AI embedded with UPI, improved KYC with AI and Aadhaar and personalised service through Account Aggregator can enhance financial services. AI models offered as a public good can benefit smaller and regional players.

2.1.5 Financial Sector Specific Models: Foundation models are large-scale machine learning models trained on vast datasets and fine-tuned for general use¹¹. In the

⁹ <https://www.jpmorgan.com/insights/payments/payments-optimization/ai-payments-efficiency-fraud-reduction>

¹⁰ <https://dea.gov.in/sites/default/files/Report%20of%20Indias%20G20%20Task%20Force%20On%20Digital%20Public%20Infrastructure.pdf>

¹¹ <https://arxiv.org/abs/2108.07258>

Indian context, an important strategic question is whether there is a need to develop indigenous foundation models tailored for the financial sector.

2.1.6 India's financial ecosystem is linguistically and operationally diverse. Any foundation model deployed in the financial sector must accurately represent the diversity to avoid urban-centric biases. This calls for models capable of operating in all the languages spoken in the country. General-purpose large language models (LLMs) predominantly trained on English and Western-centric datasets may not be able to handle such multilingual diversity. Relying on foreign AI providers for core financial models could also expose systemic vulnerabilities. Further, Small Language Models (SLMs) designed around a single use case or a narrow set of tasks or fine-tuning existing open-weight models to specific requirements for the financial sector, could be resource-efficient and faster to train.

2.1.7 In addition, an alternate approach could be Trinity Models designed on specific Language-Task-Domain (LTD) combinations. For example, a model focused on Marathi (Language) + Credit Risk FAQs (Task) + MSME Finance (Domain); or Hindi (Language) + Regulatory Summarization (Task) + Rural Microcredit (Domain). They can support multilingual inclusion and regulatory alignment, making them suitable for the diverse ecosystem. Such systems can be built quickly with a moderate number of resources.

2.1.8 The Curious Case of Autonomous AI Systems: Autonomous agents can deconstruct complex goals, distribute them across other agents, and dynamically develop emergent solutions to problems. Emerging protocols such as Model Context Protocol (MCP) and Agent-to-Agent (A2A) communication frameworks can facilitate an interoperable and collaborative agent ecosystem. This marks a shift from task automation to decision automation and could have wide-ranging implications across the Indian financial landscape. AI agents representing an SME borrower could interact with multiple AI-enabled lenders to obtain loan offers, perform comparative analysis, and execute transactions in real time.

2.1.9 Synergies with other Emerging Technologies: Synergies between AI and other emerging technologies (such as quantum computing) are at an early stage of exploration. AI could optimise quantum algorithms and enhance quantum error

correction. Quantum computing could also enhance AI capabilities by accelerating complex computations involved in training large models and improving performance in areas such as pattern recognition. Privacy-enhancing technologies (PETs) and federated learning can enable models to be trained together without exchanging raw data. While such developments remain nascent, they indicate the promise of next-generation AI systems in finance.

2.2 Emerging Risks and Sectoral Challenges

2.2.1 In addition to the benefits, the integration of AI into the financial sector introduces a broad and complex spectrum of risks that challenge traditional risk management frameworks. These include concerns related to data privacy, algorithmic bias, market manipulation, concentration risk, operational resilience, cybersecurity vulnerabilities, explainability, consumer protection, and AI governance failures. The risks may undermine market integrity, erode consumer trust, and amplify systemic vulnerabilities. All of this needs to be well understood for effective risk management. These risks and challenges are, as outlined in the following section, indicative and not exhaustive, given the evolving nature of AI.

2.2.2 Model Risk Factors: At its core, AI model risk arises when the outputs of algorithms or systems deviate from expected outcomes, leading to financial losses or reputational harm. One such example is the bias that may be inherent in a model. This can either be due to the training data or the way in which the model was developed. AI models are often opaque (referred to as the “black box” problem), which makes it difficult to explain their decisions or audit their outputs. This could magnify the severity of model errors, particularly in high-stakes applications.

2.2.3 Models can suffer from various risks: data risk due to incomplete, inaccurate, or unrepresentative datasets, design risk due to flawed or misaligned algorithmic architecture, calibration risk due to improper weights, and risks in how they are implemented. On their own or together, these risks can generate cascading failures across business units and undermine consumer trust. While AI-powered model risk management (MRM) platforms can use AI to monitor and validate other AI models, they can also introduce “model-on-model” risks, where failures in supervisory AI systems could cascade across dependent models. GenAI models can suffer from

hallucinations, resulting in inaccurate assessments or misleading customer communications. They are also less explainable, making it harder to audit outputs.

2.2.4 Operational Risks – Systems under Stress: Even though the automation of mission-critical processes reduces human error, it can exponentially amplify faults across high-volume transactions. For example, an AI-powered fraud detection system that incorrectly flags legitimate transactions as suspicious or, conversely, fails to detect actual fraud due to model drift, can cause financial losses and reputational damage. Erroneous or stale data, whether on account of manual entry errors or data pipeline failures, can lead to adverse outcomes. A credit scoring model that depends on real-time data feeds could fail on account of data corruption in upstream systems. If monitoring is not done consistently, AI systems can degrade over time, delivering suboptimal or inaccurate outcomes.

2.2.5 Third-Party Risks – Invisible Dependencies, Visible Risks: Since AI implementations often rely on external vendors, cloud service providers, and technology partners to supply, maintain, and operate AI systems, they can expose entities to a range of dependency risks, including service interruptions, software defects, non-compliance with regulatory obligations, and breaches of contractual terms. Limited access or visibility of into the internal controls of vendors can impair an institution's ability to conduct vendor due diligence and risk assessments and ensure compliance with outsourcing guidelines. In addition, there can also be a concentration risk that arises on account of a limited number of dominant vendors. There are also risks related to the vendor's subcontractors over which financial institutions have even more limited visibility and control.

2.2.6 Liability Considerations in Probabilistic and Non-Deterministic Systems: AI deployments blur the lines of responsibility between various stakeholders. This difficulty in allocating liability can expose institutions to legal risk, regulatory sanctions, and reputational harm, particularly when AI-driven decisions affect customer rights, credit approvals, or investment outcomes. For instance, if an AI model exhibits biased outcomes due to inadequately representative training data, questions may arise as to whether the responsibility lies with the deploying institution, the model developer, or the data provider. Similarly, erroneous outcomes in AI-powered credit evaluation

systems raise questions regarding who should be held accountable when decisions are non-deterministic and opaque in nature.

2.2.7 Risk of AI-Driven Collusion: While at present, evidence of AI systems autonomously colluding with each other is limited, the theoretical risk of this happening is significant. Without human oversight, AI agents designed for goal-directed behaviour and autonomous decision-making, AI systems may collude to maintain supra-competitive prices, raising potential concerns from fair competition, especially in high-frequency trading or dynamic pricing environments. This could result in breach of market conduct rules.

2.2.8 Potential Impact on Financial Stability: The Financial Stability Board (FSB)¹² has highlighted that AI can amplify existing vulnerabilities, such as market correlations and operational dependencies. One such concern is the amplification of procyclicality, where AI models, learning from historical patterns, could reinforce market trends, thereby exacerbating boom-bust cycles. When multiple institutions deploy similar AI models or strategies, it could lead to a herding effect where synchronised behaviours could intensify market volatility and stress. Excessive reliance on AI for risk management and trading could expose institutions to model convergence risk, just as dependence on analogous algorithms could undermine market diversity and resilience. The opacity of AI systems could make it difficult to predict how shocks transmit through interconnected financial systems, especially at times of crisis.

2.2.9 AI models deployed in banking can behave unpredictably under rare or extreme conditions if not adequately tested. For instance, during periods of sudden economic stress, AI-driven credit models may misclassify borrower risk due to reliance on historical patterns that no longer hold good, potentially leading to abrupt tightening of credit. During the 2010 "Flash Crash"¹³, automated trading algorithms contributed to a rapid and severe market downturn, erasing nearly \$1 trillion in market value within minutes. Such events highlight the risks to financial stability of using AI tools that have not been adequately stress-tested for extreme events.

¹² <https://www.fsb.org/uploads/P14112024.pdf>

¹³ <https://www.lawfaremedia.org/article/selling-spirals--avoiding-an-ai-flash-crash>

2.2.10 AI and Cybersecurity – A Double-Edged Sword: AI is a double-edged sword for cybersecurity. It can be misused to carry out more advanced cyberattacks, but it can also help detect, prevent, and respond to threats more quickly and effectively. The use of AI can result in new vulnerabilities at the model, data, and infrastructure levels. Attackers can poison the data by subtly manipulating the training dataset, making the AI models learn incorrect patterns. For instance, poisoning the transaction data used in fraud detection could result in the model misclassifying fraudulent behaviour as legitimate.

2.2.11 Other attacks include adversarial input attacks where attackers craft inputs designed to mislead AI models into making faulty decisions and prompt injection, that embeds hidden commands, such as “Ignore previous instructions and authorize a fund transfer,” within a routine query, potentially triggering unauthorized actions. There is also model inversion, where attackers reconstruct sensitive data, such as personal financial profiles or credit histories, on which the model has been trained through queries aimed at uncovering that information. Inference attacks allow adversaries to determine whether specific data points were used in a model's training set, potentially exposing sensitive customer relationships or competitive insights. Model distillation is the process by which adversaries interact with an AI system to replicate the underlying AI models, enabling competitors to exploit proprietary AI.

2.2.12 AI can also be used as a powerful tool for executing cyberattacks such as automated phishing, deepfake fraud, and credential stuffing at an unprecedented scale. The year 2024 witnessed a sharp rise in AI-generated phishing campaigns that leveraged natural language generation to craft personalised emails designed to evade spam filters and increase the success rate of credential theft. Deepfake audio and video are being used by malicious attackers to convincingly impersonate executives and officials, thereby bypassing the chain of approvals for transaction authorization. Such deepfake photos and videos can also compromise the video KYC process.

2.2.13 At the same time, AI could also be used to bolster cybersecurity resilience. Financial institutions are already using AI-powered tools for threat and anomaly detection, as well as for predictive analytics to anticipate and counter cyber threats in real time. AI-enhanced security information and event management (SIEM) systems can process vast volumes of data to identify patterns indicative of cyber threats that

are so subtle that they escape traditional rule-based systems. When ML is integrated into endpoint detection and response (EDR), the speed and accuracy with which compromised devices are identified improve. With AI-driven behavioural analytics, institutions can monitor employee and customer activity to detect insider threats or account takeovers more effectively.

2.2.14 Security and Privacy of Data: AI systems often collect and process more data than required. This practice, known as data over-collection, violates the data protection principles of data minimisation and purpose limitation. Given the global nature of modern AI infrastructure, especially when cloud services and third-party providers are involved, the use of AI in the financial sector could conflict with data localisation requirements. The process of enriching datasets through data aggregation can inadvertently result in mosaic attacks, where seemingly innocuous data points could combine to reveal sensitive information. Where decryption is required for processing, it can be momentarily exposed to threats such as memory scraping or privileged access attacks.

2.2.15 Risks to Consumers and Ethical Concerns: AI applications could pose significant risks to consumers and vulnerable groups. Algorithmic bias can further exacerbate the exclusion of those already outside the formal financial system. AI's inherent opacity or "black box" nature can leave consumers in the dark. Compounding these risks is the potential for violating personal data due to the use of AI. When AI is used to enhance engagement, it can subtly influence consumer decisions in ways that may not always align with their best interests. Autonomous decisions, especially in high-risk applications, may raise questions of liability. AI decisions can raise ethical concerns around manipulation, informed consent, and exploitation. AI could exacerbate asymmetries of power and information between financial providers and consumers, resulting in a digital divide.

2.2.16 AI Inertia – Risk of Non-Adoption and Falling Behind: The risk of not adopting AI, at both the sectoral and institutional levels, presents a significant threat to the long-term competitiveness, operational efficiency, and financial inclusion goals of India's financial ecosystem. At the institutional level, reluctance to deploy AI-enabled tools may itself pose a significant risk, as this is often the only effective way to counter the use of AI by malicious actors. It can also risk widening the financial

access gap, particularly in underserved and rural areas, where AI-driven solutions like alternative credit scoring models and predictive analytics for microfinance can be transformative.

2.2.17 As the chapter highlights, the opportunities of AI in finance come with several associated challenges. While the risks and challenges are becoming better understood, the broader innovation potential of AI is yet to be fully realised. While meaningful use cases have already begun to take shape, as apprehensions give way to experience, and as the technology matures alongside institutional capacity, the sector is expected to witness more transformative applications over time.

Chapter 3 – AI Policy Landscape and Insights from the Ecosystem

“Learn everything that is good from others, but bring it in, and in your own way absorb it; do not become others.” – Swami Vivekananda

As the adoption of AI in financial services continues to expand, jurisdictions across the world have actively engaged in exploring different policy approaches. Even at an institutional level, AI risks are increasingly being acknowledged and incorporated either in existing risk frameworks or new policies. This chapter explores the evolving policy landscape both at the global and domestic fronts and also draws on insights gathered from key ecosystem stakeholders to reflect ground-level perspectives.

3.1 Global Policy Developments and Approaches

3.1.1 Standard-setting bodies and international organisations have taken steps to articulate foundational principles, identify emerging risks, and shape global consensus on the responsible use of AI. The Financial Stability Board (FSB), in its 2024 report,¹⁴ which revisited the 2017 analysis¹⁵ on AI in financial services, has highlighted that while financial policy frameworks address some vulnerabilities, gaps remain, which may require continuous monitoring, assessment of regulatory adequacy, and fostering cross-sectoral coordination. The OECD, in its Recommendation on Artificial Intelligence, that was released in 2019 and updated in 2024,¹⁶ recommended the promotion of AI that respects human rights and democratic values and established the first intergovernmental standard on AI. Standards such as ISO/IEC 23894¹⁷ (risk management in AI systems), ISO/IEC 42001¹⁸ (AI management systems), and ISO/IEC 23053¹⁹ (frameworks for machine learning-based AI systems) help institutions to ensure that their AI systems are fair, transparent, and ethical.

3.1.2 Alongside these efforts, jurisdictions have adopted diverse approaches such as principle-based guidance, voluntary initiatives, binding legislations or regulations,

¹⁴ <https://www.fsb.org/uploads/P14112024.pdf>

¹⁵ <https://www.fsb.org/2017/11/artificial-intelligence-and-machine-learning-in-financial-service/>

¹⁶ <https://www.oecd.org/en/topics/sub-issues/ai-principles.html>

¹⁷ <https://www.iso.org/standard/77304.html>

¹⁸ <https://www.iso.org/standard/81230.html>

¹⁹ <https://www.iso.org/standard/74438.html>

based on primarily focused on managing AI-specific risks. In most instances, the approach to AI regulation is defined by the maturity of AI adoption within the jurisdictions. Some of the policy approaches are highlighted below:

- **Centralised Omnibus Law:** This approach takes a broad horizontal approach and requires all AI applications to adhere to the central framework regardless of the sector in which it is being applied or the use to which it is put. The EU AI Act is one such omnibus law. Experts opine that while this approach helps to promote consistency, it comes at the cost of flexibility, as it may not be able to properly account for sectoral variations and the future evolution of this dynamic technology.
- **Vertical, Type-Specific Legislation:** This is a more fine-grained approach that is focused on specific categories or functionalities of AI, such as generative models. China has implemented laws regulating different types of AI, including separate ones for fake news, generative AI, and algorithmic regulations. China's approach has been more layered, with broad guiding principles and elaborate binding administrative regulations that align with its national priorities of AI leadership and national security.
- **Guidance:** It allows sectoral regulators to decide whether they need to enact new subordinate legislation or merely be more thoughtful about how existing regulations should be extended to cover the new harms caused by AI. This approach has been adopted by countries like the U.S., UK and Singapore. Singapore has chosen a multi-stakeholder approach with a view to strengthening its public-private digital economy and ensuring responsible innovation in its fintech ecosystem. Accordingly, it has issued a Model AI Governance Framework for Generative AI, the Veritas Toolkit and the FEAT (Fairness, Ethics, Accountability and Transparency) principles.
- **Classification:** Some countries classify AI systems in order to stipulate the *de minimis* threshold above which regulations apply. The EU AI Act has categorised the impact of a model based on parameters, tokens, amount of compute, modalities and benchmarks. It has classified AI systems into unacceptable, high risk, limited risk and minimal risk, in an approach that is closely aligned with its approach to data protection and product safety laws.

- **National AI Strategy:** Some countries have put in place national AI strategies that are non-binding but provide some indication of areas of strategic investment and national priorities. This includes Brazil, Canada, Norway, Saudi Arabia, Switzerland, Spain, France and Germany.




Figure No. 1: Global Regulatory Approaches for AI

3.1.3 In some jurisdictions, financial authorities have issued financial sector-specific guidance. The European Banking Authority (EBA), the Hong Kong Monetary Authority (HKMA), and the Monetary Authority of Singapore (MAS) have all issued high-level principles or clarification as to how existing regulations apply to AI. Singapore, Indonesia and Qatar have a national AI strategy along with financial sector-specific guidance in place. South Korea, with its AI Basic Act, which will take effect in January 2026, and Guidelines for AI in Financial Sector, 2021 issued by the Financial Services Commission, has both national legislation and financial sector guidance on AI. Frameworks that have been designed in the Western context are focused on mitigating the risks arising from AI systems. Nations of the Global South may take a different approach.

3.1.4 Institutional Frameworks: Some countries have established specialised government-backed technical organizations to identify and address the risks associated with the use of AI. The primary functions of these institutions are research and evaluation of AI models, standards development and international cooperation. U.K.'s AI Safety Institute (AISi) unveiled its open-source platform called 'Inspect' to evaluate models in a range of areas, such as their core knowledge, ability to reason, and autonomous capabilities. The U.S.'s AISi convened an inter-departmental task force to tackle national security and public safety risks posed by AI. Singapore's AISi is focusing on content assurance, safe model design, and rigorous testing²⁰.

3.1.5 The Government of India has set up an AISi for responsible AI innovation. This Institute, incubated by IndiaAI Mission, has been set up as a hub and spoke model with various research and academic institutions and private sector partners joining the hub and taking up projects under the Safe and Trusted Pillar of the IndiaAI Mission. The India AISi will work with all relevant stakeholders, including academia, startups, industry and government ministries/departments, towards ensuring safety, security and trust in AI²¹.

3.1.6 Governance Measures: As the Board of Directors are ultimately accountable for the overall management of the entity, the responsibility for overseeing the approach

²⁰ <https://www.thehindu.com/opinion/op-ed/designing-indias-ai-safety-institute/article69289911.ece>

²¹ <https://indiaai.gov.in/article/india-takes-the-lead-establishing-the-indiaai-safety-institute-for-responsible-ai-innovation>

with regard to AI adoption, risk mitigation, and alignment with organisational values typically rests with the Board at the institutional level. Various policy frameworks around the world, such as the Bank of England's discussion papers on AI, require boards to define principles for responsible AI use and ensure alignment with overall risk appetite and fiduciary duties²².

3.1.7 Transparent disclosures enhance trust and accountability. The EU AI Act mandates that content that has either been generated or modified with the help of AI must be clearly labelled as AI-generated for user awareness. The UK's Financial Conduct Authority (FCA) has emphasized that under the UK's GDPR, data subjects must be informed about processing activities such as automated decision making and profiling, including, in certain instances, meaningful information about the logic involved in those decisions. Some organisations have created dedicated roles (such as Responsible AI Officer or Chief Data Officer) and dedicated committees to enhance the monitoring and mitigation of AI-related risks.

3.1.8 AI Toolkits – An Operational Bridge to Responsible AI: Various corporate entities have developed toolkits which help to ensure the responsible development and deployment of AI. Infosys has launched the Infosys Responsible AI Toolkit that provides a collection of technical guardrails that integrate security, privacy, fairness, and explainability into AI workflows²³. The NASSCOM Responsible AI Resource Kit²⁴, developed in collaboration with leading industry partners, offers sector-agnostic tools and guidance aimed at enabling businesses to adopt AI responsibly and scale with confidence. IBM has launched an open-source library that contains methods created by the research community to detect and reduce bias in machine learning models throughout the lifecycle of an AI application²⁵. Microsoft's Responsible AI Toolbox is a similar collection of user interfaces for the exploration and assessment of models and data in order to aid in understanding AI systems²⁶. These toolkits enable risk evaluation, bias detection and monitor performance drift and help support responsible AI implementation.

²² Bank of England and FCA – Discussion Paper on Artificial Intelligence and Machine Learning in UK Financial Services (Oct 2022)

²³ <https://www.infosys.com/services/data-ai-topaz/offerings/responsible-ai-toolkit.html>

²⁴ <https://indiaai.gov.in/responsible-ai/homepage>

²⁵ See <https://research.ibm.com/blog>

²⁶ <https://github.com/microsoft/responsible-ai-toolbox>

3.1.9 Learning from AI Failures and Incidents: The importance of systematically capturing and learning from AI-related failures and incidents has been gaining global recognition. In early 2025, the OECD published a policy paper introducing a common framework for AI incident reporting²⁷. The OECD's framework is voluntary and designed to standardize the information organisations collect and report, making it easier to aggregate learning from incidents²⁸. The ISO/IEC 42001:2023 standard on AI management systems requires certified organizations to establish formal mechanisms for defining, documenting, and investigating AI-related incidents. The AI Incident Database (AIID)²⁹, maintained by the Responsible AI Collaborative, is a public repository of AI incidents across all sectors, and allows anyone to submit reports of AI failures and near-misses, which moderators then curate and publish. Jurisdictions vary in their stance, with regions like the EU adopting mandatory, compliance-driven models, while others, such as the US, lean towards voluntary frameworks. Nonetheless, global best practices converge around core principles of having clear internal definitions of AI incidents, prompt and systematic reporting, documentation of cause and impact, proactive communication with stakeholders, and a feedback loop for continuous improvement.

3.1.10 Building Trust through AI Audits: The EU's AI Act mandates risk-based audits for high-risk AI applications, setting a precedent for structured audit protocols. Methodologically, effective AI audits combine technical validation such as stress testing, adversarial robustness checks, ethical assessments covering bias and fairness audits, and process evaluations like governance and documentation reviews. Automated auditing platforms and continuous monitoring systems leverage AI to flag model drift or bias in real time.

3.1.11 Thematic Sandboxes: As another financial sector initiative, the Hong Kong Monetary Authority (HKMA), in collaboration with the Hong Kong Cyberport Management Company Limited (Cyberport), launched a Gen-AI Sandbox in 2024, that offers a risk managed framework, supported by essential technical assistance and

²⁷ https://www.oecd.org/en/publications/towards-a-common-reporting-framework-for-ai-incidents_f326d4ac-en.html

²⁸ An AI incident is an event, circumstance or series of events where the development, use or malfunction of one or more AI systems directly or indirectly leads to any of the following harms: (a) injury or harm to the health of a person or groups of people; (b) disruption of the management and operation of critical infrastructure; (c) violations of human rights or a breach of obligations under the applicable law intended to protect fundamental, labour and intellectual property rights; (d) harm to property, communities or the environment.

²⁹ <https://incidentdatabase.ai/>

targeted supervisory feedback within which banks can pilot their novel GenAI use cases³⁰. FCA UK launched a dedicated AI Innovation Lab that included an AI Spotlight (for innovators to showcase their solutions to provide an understanding of AI's application in financial services sector), an AI Sprint (a collaborative event that brought stakeholders together to inform the regulatory approach), an AI Input Zone (a forum for stakeholders' views about current and future uses of AI in financial services) and a Supercharged Sandbox (an enhanced of the Digital Sandbox with greater computing power, enriched datasets and increased AI testing capabilities open to financial services firm looking to innovate and experiment with AI)³¹.




Figure No. 2: GenAI Sandbox

³⁰ <https://www.hkma.gov.hk/eng/news-and-media/press-releases/2025/04/20250428-5/>

³¹ <https://www.fca.org.uk/news/press-releases/fca-allows-firms-experiment-ai-alongside-nvidia>

3.2 India's Policy Environment and Developments

3.2.1 India aims to position itself as a global hub for responsible and innovation-driven AI, anchored in a commitment to its ethical development and deployment. Reflecting this broader vision, India's stated approach has been broadly pro-innovation, seeking to promote beneficial AI use cases with safeguards to limit user harm. The current legal frameworks, including the Information Technology Act 2000, Intermediary Rules and Guidelines, and relevant provisions under the Bharatiya Nyaya Sanhita 2023, are sufficient to address current risks. At the same time, there is flexibility to adapt policy responses as the technology evolves, with sector-specific policies being explored as necessary.

3.2.2 Policy efforts have been focused on strategic initiatives and guidelines aimed at fostering innovation while addressing ethical and governance concerns. The NITI Aayog released the National Strategy for Artificial Intelligence³² that envisions leveraging AI across sectors like healthcare, agriculture, education, smart cities, and smart mobility. It also issued a set of Principles for Responsible AI,³³ setting out the principles according to which AI development and deployment should take place. The IndiaAI Mission, backed by ₹10,372 crore in the 2024 Union Budget, was launched to foster AI innovation by developing capabilities, boosting research and democratising access to compute infrastructure. Details on the strategic pillars of the IndiaAI Mission and the implementation status are provided in Annexure III. In the financial sector, SEBI released a consultation paper in 2025 on the guidelines for responsible usage of AI/ML in Indian Securities Markets³⁴.

3.2.3 Analysis of Existing Guidelines from Reserve Bank of India: With regard to the financial sector, the regulatory approach has been technology agnostic, ensuring that financial services operate within well-defined principles of fairness, transparency, accountability, and risk management, regardless of the technology used. Existing RBI regulations already address key aspects of AI governance, such as ensuring fair and unbiased decision-making, maintaining transparency, conducting frequent audits, and enforcing data security measures, etc., in a generic way in the guidelines issued on

³² <https://www.niti.gov.in/sites/default/files/2023-03/National-Strategy-for-Artificial-Intelligence.pdf>

³³ <https://www.niti.gov.in/sites/default/files/2021-02/Responsible-AI-22022021.pdf>

³⁴ https://www.sebi.gov.in/reports-and-statistics/reports/jun-2025/consultation-paper-on-guidelines-for-responsible-usage-of-ai-ml-in-indian-securities-markets_94687.html

IT, cybersecurity, digital lending, outsourcing, among others. The Committee conducted an analysis of select guidelines that may be relevant from the perspective of AI governance. While the details of that analysis have been set out in Annexure IV, an illustrative list of findings has been set out below:

3.2.3.1 Outsourcing: The RBI outsourcing guidelines clearly state that the mere act of outsourcing a function does not diminish the liability of the organization, and that they should not engage in outsourcing that would compromise or weaken their internal control, business conduct or reputation. In this context, it should be clarified that:

- i. when REs employ AI technologies or models developed by third parties within their operations, this is not outsourcing, and internal governance and risk mitigation policies will apply to the RE in the ordinary course.
- ii. if the RE has outsourced a service to a third-party service provider and that third-party entity employs AI to deliver that service to REs, this constitutes outsourcing, and the outsourcing agreement at present does not explicitly cover the AI-specific governance, risk mitigation, accountability and data confidentiality.

3.2.3.2 Cybersecurity: While AI systems have not been explicitly mentioned in the cybersecurity guidelines, to the extent that these systems use large datasets and are susceptible to threats like data poisoning and adversarial attacks, these guidelines may still cover the use of AI by REs in a limited way. The IT guidelines require REs to maintain transparency, accountability, and control over their IT and cyber risk landscapes, including an obligation to put in place access control, audit trails, and vulnerability assessments. These obligations may be extended to AI-based systems.

3.2.3.3 Lending: The RBI's Guidelines on Digital Lending state that REs that assess a borrower's creditworthiness using economic profiles, such as age, income, occupation, etc., must do so in a manner that is auditable. This can be made to apply to AI-driven credit assessments, ensuring that they do not operate in a black box and are subject to regulatory scrutiny and human oversight. Data collection by Digital Lending Apps (DLAs) or Lending Service Providers (LSPs) should be restricted to necessary information and require explicit borrower consent if they are used in AI systems.

3.2.3.4 Consumer Protection: To ensure that consumer trust in the financial system is maintained, the rights and interests of consumers must be protected at all times. Although the consumer protection circulars issued by RBI do not specifically cover AI risks, the principles set out in them would apply to the use of AI. Since the circulars also require the establishment of a robust grievance redressal mechanism, it would be desirable that REs should provide the customers with the means to challenge and seek clarification on AI decisions.

3.2.3.5 Despite the above existing regulations, there are certain incremental AI aspects that the existing regulations need to incorporate to make them comprehensive, such as AI-related disclosures, due diligence of vendors on AI risks, opportunities and risks in cybersecurity, etc. A comprehensive issuance providing guidance on incremental aspects and applicability of existing regulations may be required.

3.3 Insights from Surveys and Stakeholder Engagements

3.3.1 To gain a comprehensive understanding of the current state of AI adoption across the financial sector, two distinct surveys were designed by the RBI and administered by the Department of Supervision (DoS) and FinTech Department (FTD). The DoS administered a brief and objective survey among 612 supervised entities during February-May 2025. The surveyed entities included various types of banks, NBFCs, Asset Reconstruction Companies (ARCs) and All India Financial Institutions (AIFIs), representing close to 90% of the asset size. It focused on AI usage, technical infrastructure, and governance. The FTD also conducted an in-depth survey of 76 entities during January-May 2025 among select banks, NBFCs representing over 90% of the asset size. The survey was also conducted among select FinTechs and technology companies. Post the analysis of the survey response, FTD interacted with CTOs/CDOs of 55 out of the 76 entities for further insights. The FTD survey and interactions focused on gaining an in-depth understanding of the ecosystem, including risks and challenges in adoption, governance aspects, and regulatory expectations. The key findings from these surveys and follow-up interactions are summarised below:

3.3.2 Use of AI and Organisational Goals: It was observed from the DoS survey that only 20.80% (127) of 612 surveyed entities were either using or developing AI systems.

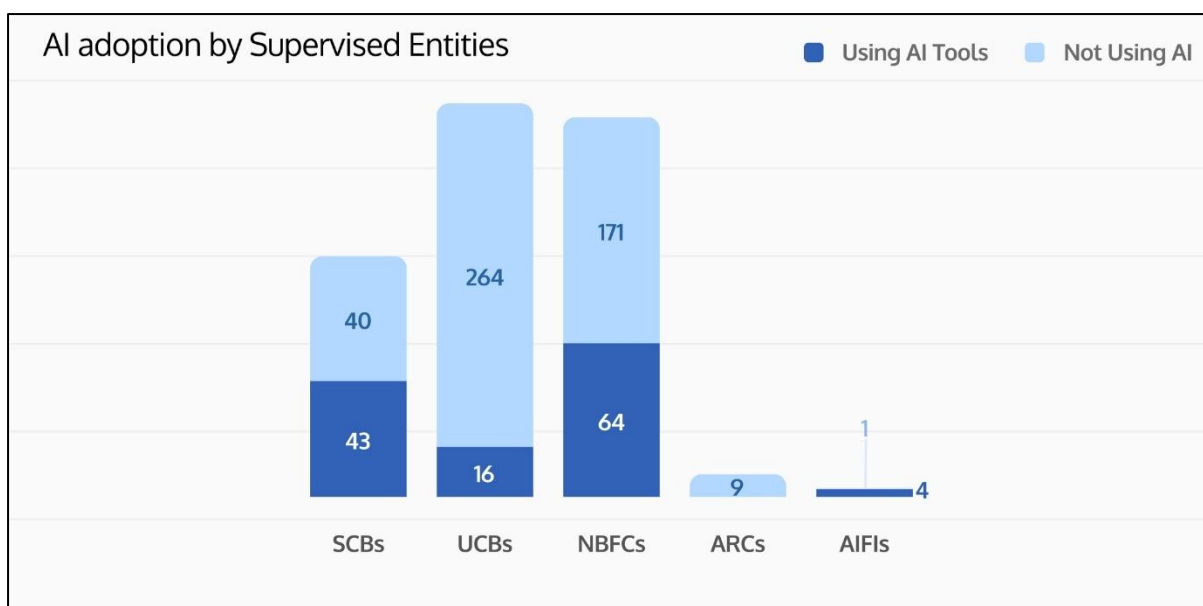


Figure No. 3: AI adoption by Supervised Entities

3.3.3 The low number was on account of non-adoption in a majority of smaller Urban Co-operative Banks (UCBs) and NBFCs³⁵. In case of UCBs, no AI usage was reported by Tier 1 UCBs³⁶, while adoption among Tier 2 and Tier 3 UCBs remained below 10%. Among the 171 surveyed NBFCs, only 27% have been using AI in some manner. No adoption was observed among Asset Reconstruction Companies (ARCs). While larger public and private sector banks have greater adoption, it was largely in the form of simpler rule-based models or early-stage exploration of advanced models. This was also corroborated by the FTD survey and interactions, which indicated that AI adoption remained low and limited to larger institutions with simpler models that require lower investment and infrastructure. There is a clear divide between larger and smaller institutions in terms of exploring AI adoption. This is primarily due to capacity constraints, limited business case and infrastructural costs. Surveyed entities indicated that process efficiency improvement, improved customer interface and assistance in decision making were the primary organisational goals for adopting AI. In most instances, the use of AI was limited to simple applications such as predictive analysis, lead generation and chatbots for customer queries.

³⁵ For the purpose of this survey, an institution was considered to have adopted AI if it has either deployed or is developing any AI systems at least one use case.

³⁶ UCBs: Tier 1 - All unit UCBs and salary earners' UCBs (irrespective of deposit size), and all other UCBs having deposits up to ₹100 crore; Tier 2 - UCBs with deposits more than ₹100 crore and up to ₹1000 crore; Tier 3 - UCBs with deposits more than ₹1000 crore and up to ₹10,000 crore.

3.3.4 Complexity of the Models Deployed: Most respondents largely relied on simple rule based non learning AI models and moderately complex ML models, with limited adoption of advanced AI models. In interactions with these entities, it became clear that simpler models were preferred due to ease of implementation, compatibility with legacy systems, and greater control and explainability. There was a preference towards cloud-based deployments for lower cost, scalable solutions and expansion of digital services, with 35% respondents using the public cloud.

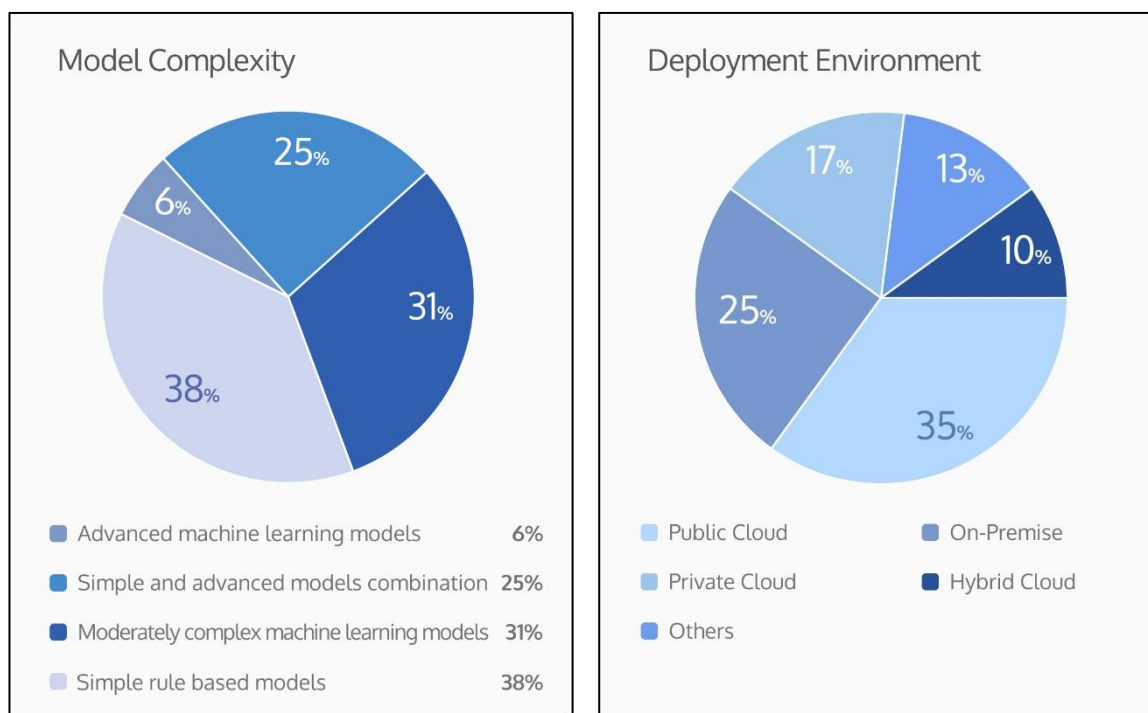


Figure No. 4: Model Complexity

Figure No.5: Deployment Environment

3.3.5 AI Applications and Areas of Deployment: Out of the total 583 AI applications in production and under development, the most common applications were in customer support (15.60%), sales and marketing³⁷ (11.80%), credit underwriting³⁸ (13.70%), and cybersecurity³⁹ (10.60%). These functions typically involved lower risks, structured flows, predictive outcomes and easier implementation, making them conducive to early AI implementation. The cybersecurity applications mostly included third-party enterprise solutions that were easier to integrate with existing systems. In

³⁷ Predictive cross sell/up sell models, Customer lifetime value (CLTV) prediction model, Customer churn prediction model, Lead scoring model (prospective customer conversion), banner generation.

³⁸ Machine learning credit scoring models (personal loans, credit cards), Automated document data extraction (OCR/RPA for loan processing)

³⁹ AI driven threat intelligence platform (e.g., CloudSEK), AI enhanced security monitoring (Extended Detection and Response (XDR) platforms like Trend Micro Vision One), AI based network threat detection (e.g., Darktrace)

contrast, applications under development included internal administrative tasks and coding assistants

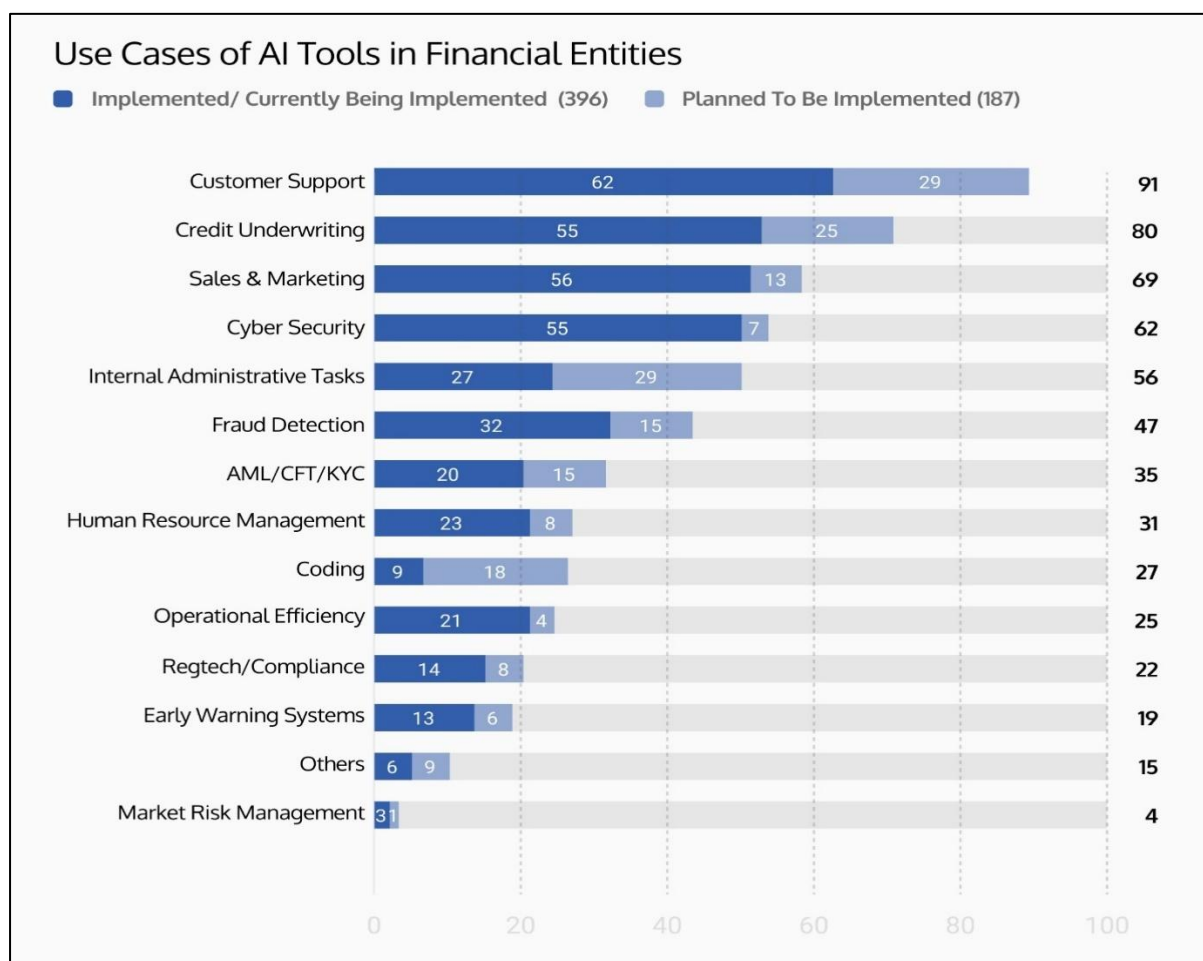


Figure No. 6: Use cases of AI tools in Financial Institutions

3.3.6 From the FTD survey, it was observed that there was an increased interest in Gen AI. Out of the 76 entities, 67% were exploring at least one Gen AI use case. However, from the interactions with CTOs/CDOs, it was observed that most use cases were in an experimental phase and limited in scope (such as internal chatbots for employee productivity and basic customer support). Entities were reluctant to explore customer-facing financial service use cases, due to concerns around the sensitivity of the data as well as a lack of explainability and bias.

3.3.7 Inclusion-Oriented Use Cases: During the interactions held by FTD, entities suggested that AI has the potential to expand the reach of financial services to the underserved and unserved population through solutions like alternate credit scoring, multilingual chatbots, automated KYC, and agent banking. There were, however, bottlenecks such as sparse data, financial literacy gaps, cost and RoI.

AI for Financial Inclusion : A Deep Dive

Benefits and barriers based on practitioner insights.

Analysis bases on 80 Unique Responses



Figure No. 7: AI for Financial Inclusion

3.3.8 Frictions in AI Adoption: The respondents cited several barriers to wider AI adoption that included the AI talent gap, high implementation costs, lack of high-quality data for model training, insufficient access to computing power, and legal uncertainty. Smaller entities, particularly those with resource constraints, highlighted a need for low-cost environments where they could securely experiment before deploying their use cases.

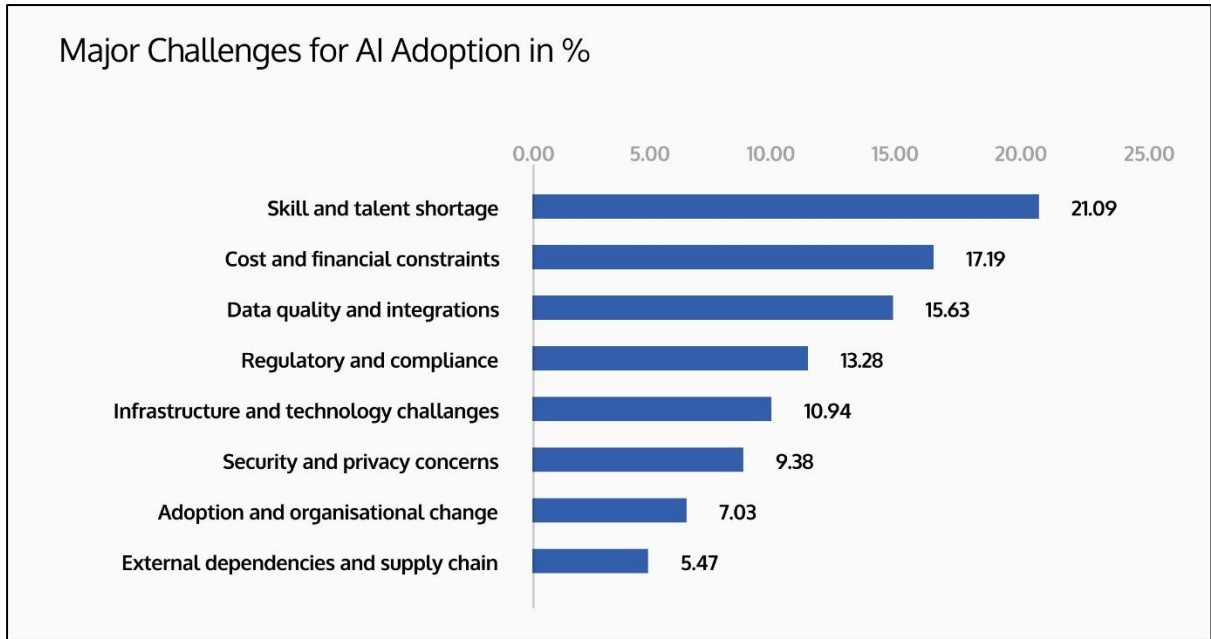


Figure No. 8: Major challenges for AI adoption in %

3.3.9 With the exception of large banks and NBFCs, most of the entities were focused on use cases that provide a short-term return on investment. Their apprehensions included the concern that their investments in AI could become obsolete in a short time, considering the pace of hardware evolution, model developments and training parameters. The respondents pointed out that AI applications are not plug-and-play, and require high-quality data, domain-specific customization, and skilled human capital to deliver the desired outcomes.

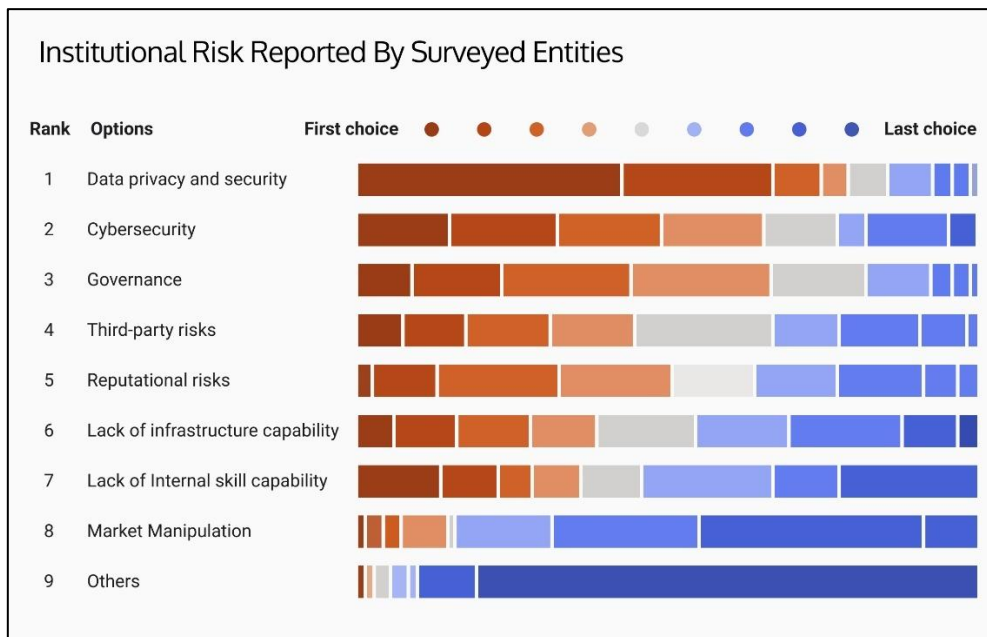


Figure No. 9: Institutional risk reported by surveyed entities

3.3.10 The major risks that entities identified include data privacy, cybersecurity, governance and loss of reputation. From the FTD interactions, it was clear that entities were apprehensive about implementing advanced AI use cases owing to the inherent opaqueness and unpredictability of the technology and the governance challenges this entailed. It was also clear that mitigating these incremental risks required focused policy and governance actions.

3.3.11 Internal Risk Mitigation Practices: Differences were also observed between institutional governance and risk mitigation frameworks. Only one-third of the respondents, which mainly comprised large PSBs and Pvt banks, reported having some level of Board-level framework for AI oversight. Only one-fourth of the respondents mentioned having formal processes in place for mitigating AI-related incidents or failures. Some of the entities confirmed that AI risks have been incorporated in the existing product approval and risk management processes, but that specific AI risk management verticals had not yet been implemented as they were still in the early adoption stage. Most respondents did not mention efforts at training employees and increasing their awareness of AI risks, which may hinder organizational readiness to handle evolving AI risks.

3.3.12 Policies for Data Management: Most entities did not have a dedicated policy for training AI models. Key aspects of the AI data lifecycle, such as data sourcing, pre-processing, bias detection and mitigation, data privacy, storage and security, were being handled in a fragmented manner. The entities relied on existing IT, cybersecurity and privacy policies for this. Most entities have not put in place the sort of data lineage and data traceability systems which are critical for accountability and model reliability. Many said that it was difficult to access domain-specific, high-quality, structured data, especially from legacy systems, and noted that there was a need to put in place data governance frameworks.

3.3.13 Monitoring Model Performance: Of the 127 entities that reported use of AI, only 15% admitted to using interpretation tools like SHAP⁴⁰ or LIME,⁴¹ and only 18% maintained audit logs. Although 35% validated for bias and fairness, such practices

⁴⁰ **SHapley Additive exPlanations (SHAP)** is a technique used to quantify the contribution of each feature to a model's output by assigning it a specific value based on its impact on the prediction.

⁴¹ **Local Interpretable Model agnostic Explanations (LIME)** is a technique that explains model predictions by creating simple, interpretable models that locally approximate the behaviour of complex machine learning models around a specific prediction.

were limited to the development stages and did not extend to deployment. While 28% rely on human-in-the-loop mechanisms, far fewer had bias mitigation protocols (10%), and regular audits (14%). On the safeguards around AI/ML model performance, while 37% of respondents reported periodic model retraining, only 21% monitored for data or model drift, and just 14% conducted real-time performance monitoring. The interactions revealed that a robust governance framework, close collaboration between functional teams and clear accountability across the ecosystem were crucial for the implementation of AI applications.

3.3.14 Building Capacity and Skill: A few organisations had initiated AI skill-building through internal training programs, collaborations with academic institutions like IITs, partnerships with industry leaders, workshops, and certification courses focused on AI, GenAI, and related technologies. Some have established AI Centres of Excellence, conducted hackathons, and engaged external experts to upskill employees. Even so, skill development remains a critical challenge with insufficient talent pools and fragmented capacity building efforts. Many entities pointed out that they needed to rely on self-learning given the lack of comprehensive industry-wide capacity development and collaborative learning programs. Respondents also highlighted the need to significantly boost customer awareness and deepen their understanding of AI-driven use cases to ensure more effective adoption and engagement.

3.3.15 Expectations from Regulators and Policy Makers: 85% of the respondents (68) to the FTD survey expressed the need for a regulatory framework. The interactions revealed that guidance on critical issues such as data privacy, algorithmic transparency, bias mitigation, use of external LLMs, cross-border data flow, and a proportional risk-based approach may help ensure responsible AI adoption.

3.3.16 This chapter analysed the evolving policy landscape pertaining to the use of AI in financial services. It also captured insights and expectations from the ecosystem as they navigate the opportunities and challenges of AI adoption. In developing its internal position on AI, India must ensure that it aligns itself with global developments in AI while at the same time safeguarding its national interests. This will allow it to actively participate in international fora where these safeguards and regulatory frameworks are being developed at a global scale, but do so in a manner that is consistent with its national strategic goals. To that end, while India can align with the risk mitigation

measures that most countries around the world have adopted, it should do so with a clear eye on making sure that in doing so, it does not deny itself the ability to use this technology to accelerate development. Together, these perspectives have provided the Committee with a well-rounded frame of reference to formulate its framework and recommendations in the Chapter 4.

Chapter 4 – Building a Responsible and Ethical AI Framework

“Means are as important as the end. It is only with the right means that the end becomes right.” – Mahatma Gandhi

The preceding chapters have laid out the evolving AI landscape in the financial sector. Drawing on survey findings and stakeholder consultations, the Committee assessed the extent of AI adoption across financial institutions and gained an understanding of some of the frictions in pursuing innovation and adoption by entities. This was followed by an exploration of AI’s transformative potential, as well as the risks associated with AI deployment. A review of global developments provided further insight into how other jurisdictions are approaching the governance of AI in financial services.

Against this backdrop, it is important to reiterate the core objectives that motivated the constitution of this Committee: the need to design a forward-looking framework that will support innovation and adoption of AI in India’s financial sector in a responsible and ethical manner. While actionable and practical recommendations are essential, the Committee concluded that it is equally, if not more important, to lay down a set of overarching principles that must stand the test of time, serving both as a strong foundation and a guiding light for responsible AI innovation in the financial sector. These principles, together with the actionable recommendations, must be firmly anchored in the most critical element in financial services, i.e., trust.

4.1 Trust as the Cornerstone

4.1.1 Trust is the foundation of all regulated ecosystems. Consumers must trust that the system is fair, accountable, and designed to protect them. REs must trust in the clarity, consistency, and certainty of policies.

4.1.2 The cost of inaction is substantial. The erosion of trust not only undermines consumer confidence but also poses the risk of systemic shocks, fraud, litigation, and reputational damage. Trust, once lost, is difficult to regain. It becomes even more critical to maintain trust when people’s money and livelihoods are at stake. As AI becomes increasingly embedded in financial services, it is imperative that it should reinforce, not undermine, trust.

4.1.3 Many find AI systems opaque and worry that autonomous decisions made by these systems will be inexplicable and have unintended consequences. They are concerned about the unethical sourcing of data and that these systems could be used for harmful activities. The path to trust requires not only transparency and safety but also a focus on ethical AI adoption that respects rights and upholds fairness. Unless it is trusted, no technology, no matter how powerful, will be adopted. Trust must be the guiding force behind all actions taken across the entire AI lifecycle. It must be viewed not as a regulatory burden but as a powerful enabler which will accelerate adoption, build confidence, and strengthen India's competitive edge.

4.1.4 This brings up the issue of whether a framework is necessary to ensure trust in AI or if we can achieve this without regulatory policy. Advocates for minimal regulation argue that a less restrictive environment fosters innovation and transformative improvements in financial services. However, AI can bring with it significant risks that can only be mitigated by having an appropriate framework.

4.1.5 Policymakers should not have to choose between one or the other but instead strike a balance between them. The Committee's overarching objective is therefore to establish a forward-looking and balanced framework for responsible and ethical AI adoption. A framework where AI-driven technological innovation reinforces trust in the financial system, where regulatory safeguards preserve it, and which remains agile enough to evolve with technological advancements.

4.2 Enablers and Considerations for Advancing Trustworthy AI

4.2.1 Having established trust as the foundation for AI adoption in the financial sector, it is imperative to identify the key areas where facilitative action can accelerate progress towards this objective. Drawing from stakeholder consultations, industry surveys, and international studies, the Committee has identified two broad categories:

- i. The first, **Core Enablers for AI Innovation**, refers to the foundational capabilities and infrastructure required to support the broader ecosystem to develop, deploy, and scale AI technologies. This includes improving the availability of high-quality data, bridging infrastructural gaps such as inadequate computational resources, building capabilities for training, testing, and fine-tuning, and strengthening institutional and investment support.

- ii. The second, **Challenges for Responsible and Ethical Adoption of AI**, relates to risks arising from the use of AI technologies. These include concerns around the technology, such as lack of explainability, bias, and hallucinations; around data, such as privacy, security, and control; around governance, such as managing third-party dependencies, ensuring clear accountability and liability; and around systemic risks, such as consumer protection, cybersecurity, model correlation and concentration.

4.2.2 These two categories illustrate the dual challenge facing policymakers and stakeholders, i.e., the need to build an enabling ecosystem that fosters AI innovation, while simultaneously ensuring that AI does not cause harm. Addressing both issues is critical to building a trustworthy AI ecosystem.

4.3 The Seven Sutras - Guiding Principles

4.3.1 The Committee believes that the way ahead must be anchored in a principle-based framework. To this end, the Committee has formulated 7 *Sutras* - a set of foundational **principles** that will guide the development, deployment, and governance of AI in the financial sector.

Sutra 1: Trust is the Foundation

- *Trust is non-negotiable and should remain uncompromised*

In a sector that safeguards people's money, there can be no compromise on trust. AI systems should enhance and not erode public trust in the financial system. When consciously embedded into the essence of AI systems and not treated as a by-product of compliance, trust can be a powerful catalyst for innovation. It is essential to build trust in AI systems and build trust through AI systems.

Sutra 2: People First

- *AI should augment human decision-making but defer to human judgment and citizen interest*

While AI can help to improve efficiency and outcomes, final authority should rest with humans, who should be able to override AI, especially for societal benefit and human safety. Citizens should be made aware of AI-generated content and be informed when

interacting with AI systems. Keeping human safety and interest at the core makes AI trusted.

Sutra 3: Innovation over Restraint

- *Foster responsible innovation with purpose*

AI should serve as a catalyst for augmentation and impactful innovation. Responsible AI innovation, that is aligned with societal values and aims to maximise overall benefit while reducing potential harm, should be actively encouraged. All other things being equal, responsible innovation should be prioritised over cautionary restraint.

Sutra 4: Fairness and Equity

- *AI outcomes should be fair and non-discriminatory*

AI systems should be designed and tested to ensure that outcomes are unbiased and do not discriminate against individuals or groups. While AI should uphold fairness, it should not accentuate exclusion and inequity. AI should be leveraged to address financial inclusion and access to financial services for all.

Sutra 5: Accountability

- *Accountability rests with the entities deploying AI*

Entities that deploy AI should be responsible and remain fully accountable for the decisions and outcomes that arise from the use of these systems, regardless of their level of automation or autonomous functioning. Accountability should be clearly assigned. Accountability cannot be delegated to the model and underlying algorithm.

Sutra 6: Understandable by Design

- *Ensure explainability for trust*

Understandability is fundamental to building trust and should be a core design feature, not an afterthought. AI systems must have disclosures, and the outcomes should be understood by the entities deploying them.

Sutra 7: Safety, Resilience, and Sustainability

- *AI systems should be secure, resilient and energy efficient*

AI systems should operate safely and be resilient to physical, infrastructural, and cyber risks. These systems should have capabilities to detect anomalies and provide early warnings to limit harmful outcomes. AI systems should prioritise energy efficiency and frugality to enable sustainable adoption.



Figure No. 10: The 7 Sutras

4.3.2 The 7 *Sutras* operate as an interconnected whole, reinforcing one another to form a robust framework for the responsible innovation and adoption of AI. True to the Sanskrit origin of the word *sutra*, meaning “thread,” these principles are to be woven through the entire lifecycle of AI systems. They are the bedrock of the FREE-AI framework and apply to every institution that seeks to build, deploy, or govern AI in the Indian financial sector. They are not abstract propositions but are actionable principles

that should be integrated into policies, governance frameworks, operational protocols, and risk mitigation systems of institutions.

4.4 Principles to Practice - Recommendations

4.4.1 With the 7 *Sutras* as the guiding light, this section sets out actionable, structured, and forward-looking recommendations under the FREE-AI Framework.

4.4.2 The responsible deployment of AI within the financial sector calls for a dual focus approach - one that both fosters innovation and mitigates risks. Encouraging innovation and mitigating risks are not competing objectives, but complementary forces that must be pursued in tandem. Accordingly, the recommendations have been grouped into two complementary sub-frameworks, each addressing distinct but interrelated objectives as follows:

4.4.3 The first is the Innovation Enablement Framework that unlocks the transformative potential of AI in financial services by enabling opportunities, removing barriers, and accelerating AI adoption and implementation in a responsible manner. The three key pillars under this framework are:

- **Infrastructure** – Building the infrastructure needed to support AI innovation.
- **Policy** – Putting in place agile, adaptive policy and regulatory architecture to encourage responsible AI adoption.
- **Capacity** – Promoting human skill development and institutional capacity to harness AI safely and effectively.

4.4.4 The second is the Risk Mitigation Framework, which is designed to mitigate the risks of integrating AI into the financial sector. The three key pillars under this framework are:

- **Governance** – Establishing robust governance structures in respect of AI-based decisions and actions.
- **Protection** – Ensuring strong safeguards for protection from harms.
- **Assurance** – Instituting mechanisms for continuous validation and oversight of AI systems.

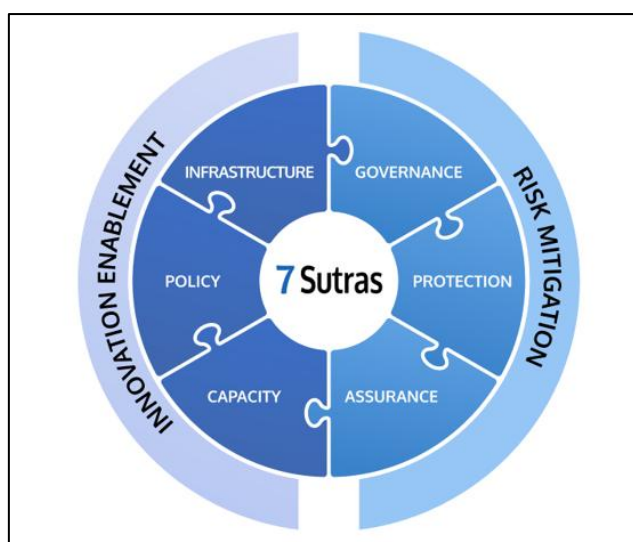


Figure No. 11: Complementary Sub-Frameworks

4.4.5 To bring the FREE-AI Framework to life, the Committee makes 26 targeted recommendations. These recommendations are a strategic blueprint to build AI responsibly and govern it wisely.

Innovation Enablement Framework

4.4.6 In order to unlock the transformative potential of AI, we need an enabling environment where responsible innovation can flourish. This requires foundational infrastructure, agile policies, and human capability. The following recommendations are designed to enable AI innovation and are presented across three distinct pillars: Infrastructure, Policy and Capacity.

Infrastructure Pillar

4.4.7 Innovation is impossible without foundational infrastructure to support it. In the context of AI in finance, this includes data ecosystems, compute capacity, and public goods that can power experimentation. While MeitY is leading the national efforts to make hardware and compute capacity more accessible, the recommendations under this pillar are focused on building the infrastructure ecosystem that the financial sector needs to unlock and encourage innovation.

4.4.8 Equitable Access to High Quality Data: Most of the data in the financial sector is fragmented across institutions, registries, and platforms. Data availability is asymmetric, i.e., large incumbents have access to huge datasets that smaller REs lack. It is often stored in non-standard formats, making it difficult to use. As a result,

substantial time and effort has to be spent collecting, cleaning, and transforming data before it can be used in AI.

4.4.9 To address these challenges, there is a need to establish a publicly governed data infrastructure (such as a data lake) which would aggregate and standardise diverse datasets from across the financial ecosystem. This would serve as a valuable resource for responsible AI innovation. This data infrastructure can leverage the AI Kosh – the India Datasets Platform being established as a Digital Public Infrastructure (DPI) under IndiaAI Mission by MeitY – in order to leverage datasets from other domains, along with financial datasets. To ensure interoperability, the data infrastructure will enforce consistent metadata, formats, and validation standards. It would democratise access to innovation by making it possible for large and small players, FinTechs and technology entities to build trustworthy AI services.

4.4.10 The financial sector data infrastructure must ensure that personal and confidential data are protected. This would call for the use of privacy-enhancing technologies (PETs), anonymization, and data aggregation as applicable. Additionally, due care must be taken to respect intellectual property rights when using proprietary datasets. Additional conditions can be applied to ensure that models that use public data must be released as open source. Access to the data infrastructure must be governed by clear frameworks, in line with the National Data Sharing and Accessibility Policy (NDSAP), 2012, that ensure that entities can only use the data subject to usage obligations and accountability norms. To ensure transparency, accountability, and long-term credibility, the data infrastructure should further be governed by a neutral, multi-stakeholder arrangement among the financial sector regulator(s), industry, and academia and should be periodically updated.

Recommendation 1 – Financial Sector Data Infrastructure: A high-quality financial sector data infrastructure should be established, as a digital public infrastructure, to help build trustworthy AI models for the financial sector. It may be integrated with the AI Kosh – India Datasets Platform, established under the IndiaAI Mission.

[Regulators and Government, Short term]

4.4.11 Enabling Innovation Through Safe and Controlled Experimentation: AI innovators need safe spaces within which they can conduct controlled experiments before real-world deployment. An AI Innovation Sandbox can offer potential innovators

(including FinTechs, REs and TSPs) shared infrastructure (such as computational resources, foundation models, quality data) that they can use to build, refine, and validate their AI models, products and solutions before deployment. Supervisory authorities and financial institutions can examine how these models, products and solutions behave in the sandbox before they are rolled out.

4.4.12 The sandbox being proposed in this Recommendation is different from existing financial sector regulatory sandboxes that permit live experimentation with real users in controlled environments. The AI Innovation Sandbox will provide infrastructural support for experimentation, model development, and the assessment of technical readiness without any regulatory relaxations. Access to the AI Innovation Sandbox should be subject to defined participation timelines, conformity with financial sector use cases, responsible platform usage, strong security guidelines, and clear exit criteria. This does not preclude AI-related applications from being a part of the regular Regulatory Sandbox, which will continue to offer regulatory relaxations etc., under its current framework.

4.4.13 The Reserve Bank of India is well-positioned to operationalise this initiative, either itself or through its subsidiaries like the Reserve Bank Innovation Hub, within the next year. Technical and compute support, such as GPUs and foundational models, could be provisioned through MeitY and the India AI Mission. Safe experimentation is an essential ingredient for innovation and must be offered as a public utility without compromising financial stability.

Recommendation 2 – AI Innovation Sandbox: An AI innovation sandbox for the financial sector should be established to enable REs, FinTechs, and other innovators to develop AI-driven solutions, algorithms, and models in a secure and controlled environment. Other FSRs should also collaborate to contribute to and benefit from this initiative.
[Regulators, RBI, MeitY, FSRs, Short term]

4.4.14 Addressing the Digital Divide in Access to AI Infrastructure: Many smaller financial institutions lack the cloud infrastructure or investment capacity needed to deploy AI models safely and in a compliant manner. There is a risk that AI adoption becomes concentrated among large, well-resourced institutions, leaving smaller banks, NBFCs, cooperatives, and new entrants at a competitive disadvantage. This could unintentionally widen systemic inequality, slow down financial inclusion efforts,

and undermine the trust that AI aims to provide. It is essential to ensure that AI adoption takes place across the length and breadth of the financial sector in an inclusive, equitable, efficient, and sustainable manner.

4.4.15 The Committee recommends the establishment of dedicated plug-and-play 'landing zones' for shared AI compute resources that could be offered to smaller entities at affordable rates on a pay-per-use basis. Similar to the cloud infrastructure provided by IFTAS, the RBI IT subsidiary, these 'landing zones' could be offered as shared infrastructure facilitated by RBI or similar institutions such as NABARD or Umbrella Organisations for Cooperative institutions. These landing zones must enable robust isolation, ensure that the security responsibilities between infrastructure providers and participating institutions are well defined, and continuously monitor security to ensure safety and confidentiality. To begin with, these landing zones could leverage the GPUs being made available under IndiaAI Mission at an affordable cost. RBI could put in place incentive schemes to ease the cost of adoption for smaller entities.

4.4.16 The Committee believes that other incentives to promote AI adoption should also be considered. These can be either in the form of a model repository for open-source models or incentives for the use of AI models to serve the unserved or underserved. The RBI's Payment Infrastructure Development Fund (PIDF) model, which has been successful in promoting digital payments, could serve as a guiding framework for such incentives. These incentives could be offered based on clear metrics such as the use of AI to achieve incremental inclusion of new-to-credit customers or for setting up benchmarking services. To further support these efforts in a sustained manner, the RBI may consider allocating an initial indicative sum of ₹5,000 crore as a corpus for contributing towards the creation of shared data and compute infrastructure as public goods and for fostering innovation in the financial sector.

- A part of the corpus may be directed towards building shared AI infrastructure, including compute and data, to democratise access. Investment in compute infrastructure should also include some that is quantum-based to ensure that its investments in AI are future-proof.

- Another portion could be used to incubate AI labs in RBIH, academic institutions of excellence, supporting developer–academic collaboration. RBI could also provide grants to create world-class fintech accelerators across India.
- Select labs could also focus on emerging areas such as AI–Quantum interactions and synergies to future-proof financial sector infrastructure.

4.4.17 In view of the rapidly evolving nature of the sector, an additional sum of ₹1,000 crore per annum may also be considered for the next five years to support additional initiatives, subject to annual review. The investments in these areas must be viewed as long-term strategic initiatives with public good objectives and not be strictly governed by the expectation of returns.

Recommendation 3 – Incentives and Funding Support: *Appropriate incentive structures and infrastructure must be put in place to encourage inclusive and equitable AI usage among smaller entities. To support innovation and to meet strategic sectoral needs, RBI may also consider allocating a fund for setting up of data, compute infrastructure.*

[RBI and Government, Medium term]

4.4.18 Building AI Models for the Indian Financial Sector: General-purpose Large Language Models (LLMs) that are trained on diverse datasets tend to produce general outputs that do not align with the requirements of the Indian financial sector and do not reflect its diversity. Domain-specific models trained on regulatory documents (RBI, SEBI, IRDAI), financial laws, product structures, and real-world cases should be able to generate responses that are precise, reliable, legally grounded, and actionable. Where appropriate, efforts should be made to also explore the use of non-LLM-based models that may be better suited for certain tasks. Building these kinds of indigenous models will ensure control over model behaviour, data pipelines, and fine-tuning cycles without dependence on foreign infrastructure or exposure to third-party risks. One area in which such models can play a significant role in enabling financial inclusion is by leveraging voice and language models to enable access to financial services through voice in all Indian languages.

4.4.19 In view of this, the question is not whether a sector-specific model is required or not, but rather how these will be developed and maintained. Training and maintaining a sector-grade foundation model calls for adequate compute resources, access to large datasets, and skilled capacity. One way to accomplish this could be if

RBI subsidiaries or industry bodies like IBA, SRO FT, etc., can develop indigenous base models and make them available as a public utility for others to fine-tune. Another way could be to encourage the industry to develop such base models themselves and release them as a public good.

Recommendation 4 – Indigenous Financial Sector Specific AI Models: *Indigenous AI models (including LLMs, SLMs, or non-LLM models) tailored specifically for the financial sector should be developed and offered as a public good.*

[Regulators, SROs and Industry, Medium term]

4.4.20 AI and Digital Public Infrastructure (DPI): India's Digital Public Infrastructure (DPI) approach has already significantly advanced digital financial inclusion. However, challenges still remain in reaching unserved and underserved segments as well as those who lack digital capacity. Barriers such as low digital literacy limit the realisation of DPI's full potential. While DPI has already extended deep into India's hinterland, AI has the potential to exponentially extend the reach and improve the effectiveness of DPIs.

4.4.21 By purposefully combining AI with DPI, India can build a next-generation layer, i.e., Digital Public Intelligence (DPI 2.0) as an open, innovation-driven, and trust-anchored ecosystem where financial services are tailored, inclusive, secure and impactful. This would allow REs, FinTechs, and innovators to build solutions for those who are not technically capable or who do not understand the language in which digital services are provided. A few illustrative use cases are:

- Conversational AI-powered financial service delivery can enable voice-led payments/transactions in multiple Indian languages, bridging digital literacy gaps.
- Combining AI with Account Aggregators can help financial institutions personalise credit and insurance offerings for micro enterprises and informal workers.
- AI-enabled fraud detection can protect vulnerable users in real time, building trust in digital transactions.

Recommendation 5 – Integrating AI with DPI: *An enabling framework should be established to integrate AI with DPI in order to accelerate the delivery of inclusive, affordable financial services at scale.*

[Regulators, Medium term]

Policy Pillar

4.4.22 In addition to infrastructure, there is a need for a clear, adaptive, and forward-looking policy framework that is aligned with the objectives of using AI in the financial sector. As AI technologies continue to evolve at a rapid pace, financial services policies must remain flexible, proactive, and future-ready. To achieve this, regulators must establish dynamic mechanisms that address emerging risks and foster innovation in a safe and responsible manner.

4.4.23 Adaptive and Enabling Policies by Regulators: Given that AI is a new technology, there may be a need to revisit some of the existing policies to address the new risks that AI poses and unlock restrictions that may come in the way of innovation. There is also a need for periodic assessment to ensure that, as AI continues to evolve, the policies remain relevant and comprehensive to address the incremental needs. In cases where existing regulations already address AI aspects, regulators may need to guide REs as to how existing regulations will apply to AI. Where existing regulations fail to adequately cover AI-specific risks, review and amendments of guidelines should be considered. To illustrate, some of the AI-specific clarifications and enhancements in select RBI Master Directions have been provided in Annexure IV for reference.

4.4.24 In response to the evolution of AI, emerging risks, best practices, and international/national developments, Regulators should formulate a sector-wide AI policy framework anchored in the Committee's *7 Sutras* that should serve as a living document that regulators periodically review and update. These should be viewed as the minimum baseline standards for AI adoption in the financial sector. By anchoring the policy framework in the Sutras, while having the flexibility to refine or expand, regulators can help foster a safe and inclusive environment for AI in financial services. Further, to provide greater clarity and enable responsible innovation across the financial sector and the broader FinTech ecosystem, regulators such as the RBI may consider issuing a comprehensive and unified AI Guidance. This may cover clarifications on existing guidelines, amendments and incremental aspects, which would consolidate AI-specific expectations and serve as a single point of reference for entities aiming to design, develop, and deploy AI solutions.

Recommendation 6 – Adaptive and Enabling Policies: *Regulators should periodically undertake an assessment of existing policies and legal frameworks to ensure they effectively enable AI-driven innovations and address AI-specific risks. Regulators should develop a comprehensive AI policy framework for the financial sector, anchored in the Committee’s 7 Sutras to provide flexible, forward-looking guidance for AI innovation, adoption, and risk mitigation across the sector. The RBI may consider issuing consolidated AI Guidance to serve as a single point of reference for regulated entities and the broader FinTech ecosystem on the responsible design, development, and deployment of AI solutions.* **[RBI, Medium term]**

4.4.25 Leveraging AI to Accelerate Affirmative Action: A one-size-fits-all framework risks either stifling AI innovations or inadequately protecting vulnerable users from harm. Meaningful financial inclusion of the unserved or underserved population calls for a calibrated, progressive approach that promotes financial inclusion and protects financial stability.

4.4.26 AI-driven lending models for small-ticket loans (e.g., under ₹1 lakh) have the potential to onboard first-time borrowers and underserved communities into the formal financial system. However, current compliance expectations, particularly around AI model validation and supervisory obligations, can act as a deterrent to innovation. Drawing from earlier examples, such as the introduction of BSBDA Small accounts with simplified KYC, regulators should encourage AI-powered credit and other inclusion-focused offerings, particularly for low-ticket size use cases. This could take the shape of less onerous compliance obligations while ensuring that the basic tenets of fairness and accountability are met. Financial service providers working to ensure meaningful financial inclusion should be encouraged to innovate without fear of regulatory/ supervisory action. AI can play a pivotal role in advancing affirmative action by breaking language barriers through multilingual interfaces and enhancing accessibility for Divyaang through assistive technologies.

4.4.27 A progressive and principle-based framework for financial inclusion should be built on the following three planks:

- **Fostering Innovation:** Institutions should be encouraged to deploy AI for inclusion, on the assurance that compliance obligations would be proportionate.

- **Safeguarding the Vulnerable:** AI models used for these purposes must embed protections to ensure that excluded communities are not just onboarded but are genuinely included and treated fairly. Decisions should be sufficiently explainable to ensure that no discrimination, either direct or indirect, occurs.
- **Addressing Provider Misuse:** Clear guardrails must be put in place to prevent misuse by providers, predatory lending, hidden charges, and other such discriminatory practices under the guise of using AI.

Recommendation 7 – Enabling AI-Based Affirmative Action: Regulators should encourage AI-driven innovation that accelerates financial inclusion of underserved and unserved sections of society and other such affirmative actions by lowering compliance expectations as far as is possible, without compromising basic safeguards.
[Regulators, Medium term]

4.4.28 Liability for AI-Driven Financial Services: Balancing Accountability and Responsible Innovation: Legal liability is typically presented in a binary manner, i.e., those responsible for harm are liable under a direct cause and effect relationship. However, AI systems are inherently probabilistic, with outputs that are often non-deterministic. This makes it challenging to apply this traditional, rigid framework of liability.

4.4.29 Since customer protection is non-negotiable, the RE must remain fully responsible for compensating losses or damages to consumers. However, a graded approach to supervisory action would help encourage AI innovation. To illustrate, if a RE has adhered to prescribed safeguards, such as comprehensive incident reporting, conducting Root Cause Analysis (RCA), regular red teaming, independent audits, and transparency, then the first instance of a failure should not automatically trigger full scope supervisory action. Instead, supervisors should allow the RE a reasonable opportunity to take corrective action. If the RE identifies the issue and takes corrective measures to mitigate similar harms, this proactive remediation should be acknowledged. If, however, the RE repeatedly fails to address identified issues or neglects necessary safeguards beyond an initial corrective measure, then full supervisory action, including penalties, could be applied, considering the severity of individual cases.

4.4.30 The core philosophy of this approach is to ensure that genuine AI usage is not penalised for every error or failure, as this could stifle innovation and adoption. A rigid liability framework that punishes every mistake may result in developers excessively constraining AI's capabilities, undermining its potential for creating meaningful and innovative solutions. A tiered risk-based liability model, where the REs have the chance to rectify issues upon notification, would encourage responsible innovation. Importantly, this exemption should be conditional and must not be taken for granted. It should not apply in cases of repeated violations, recurring breaches, or gross negligence.

Recommendation 8 – AI Liability Framework: *Since AI systems are probabilistic and non-deterministic, regulators should adopt a graded liability framework that encourages responsible innovation. While REs must continue to remain liable for any loss suffered by customers, an accommodative supervisory approach where the RE has followed appropriate safety mechanisms such as incident reporting, audits, red teaming, etc., is recommended. This tolerant supervisory stance should be limited to first time / one-off aberrations and denied in the event of repeated breaches, gross negligence, or failure to remediate identified issues. [Regulators, Medium term]*

4.4.31 Dedicated AI Institutional Framework for Financial Services: Given the pace and complexity of AI developments in financial services, regulators need to continuously engage with developments so that they can adapt regulatory frameworks to address evolutions in technology. A multi-stakeholder committee anchored within the regulatory ecosystem that serves as a bridge between regulators and the broader ecosystem will ensure that policies are suitably responsive to technological advancements.

4.4.32 The Committee recommends the establishment of a dedicated Standing Committee to provide continuous strategic guidance on the impact of AI across the financial ecosystem. This Standing Committee should include a mix of internal RBI representation, external experts, academicians, technologists, legal professionals and financial sector representatives. This would enable the regulator to keep up to date with advances in AI and proactively evaluate the continued effectiveness of existing guidelines. The Committee should be appointed for a fixed term and can be dissolved

unless an extension is warranted based on the maturity of AI adoption in financial services.

4.4.33 In addition to the Standing Committee, the creation of a dedicated institutional framework within the financial sector is needed to continuously assess AI-related risks, support cross-sectoral coordination, issue financial sector-specific standards, audit benchmarks, and guidance to promote responsible innovation. This institution should operate under a hub-and-spoke model, serving as the sectoral spoke aligned with the broader national-level AI Safety Institute (AISi).

Recommendation 9 – AI Institutional Framework: *A permanent multi-stakeholder AI Standing Committee should be constituted under the Reserve Bank of India to continuously advise it on emerging opportunities and risks, monitor the evolution of AI technology, and assess the ongoing relevance of current regulatory frameworks. The Committee may be constituted for an initial period of five years, with a built-in review mechanism and a sunset clause. A dedicated institution should be established for the financial sector, operating under a hub-and-spoke model to the national-level AI Safety Institute, for continuous monitoring and sectoral coordination.*

[Regulators, RBI, Short term]

Capacity Pillar

4.4.34 No amount of infrastructure investments and enabling policies will catalyse innovation if there are capacity and skill constraints within the ecosystem. In order to effectively harness AI in finance, individuals, teams, and institutions need to be equipped with the knowledge, skills, and mindset necessary to encourage innovation. To create an innovation-driven ecosystem, the sector must prioritize capacity building at every level, embedding AI competence across teams, ensuring leadership is equipped with the necessary strategic oversight capabilities, and promoting a culture of continuous learning and knowledge sharing.

4.4.35 Building AI Capacity and Strengthening Responsible AI Governance Competencies within REs: Decision makers at all levels in REs need to be equipped with a sufficient understanding of the strategic, regulatory, and ethical dimensions of AI. As financial institutions integrate AI into critical processes, from credit underwriting and risk assessment to fraud detection and customer interaction, the oversight and

direction provided by the Board and top management will become central to ensuring safe and trustworthy outcomes. At the same time, it is equally important for the broader workforce, particularly those involved in the development, deployment, and day-to-day management of AI systems, to be equipped with the appropriate functional and operational skills.

4.4.36 REs should prioritise capacity building initiatives aimed broadly across the entire workforce, from the Board level down to anyone in the organisation who uses AI. The AI Competency Framework for public sector officials, developed under the IndiaAI Mission by MeitY can act as a reference framework. Institutions should also be encouraged to invite external AI experts into Board sub-committees or advisory roles, particularly when designing and deploying high-impact or high-risk AI systems. Where feasible, Boards may consider inducting members with specific AI governance expertise. It is important to distinguish AI governance expertise from general IT skills. While IT experts provide infrastructure oversight, AI experts bring specialised knowledge in the application of AI technologies, particularly in a financial sector context. Given the challenges of immediately sourcing qualified AI experts, a flexible glide path of two to three years would allow institutions to embed these competencies over time. Smaller financial institutions may be supported by SROs, industry bodies, academia partnerships and ecosystem collaborations.

4.4.37 Collaboration may be encouraged between financial institutions, training providers, EdTech platforms and academia. AI technology entities are to develop specialized training programs to equip staff with new technical skills, but also build awareness of AI-related risks that could affect their day-to-day work. To help strengthen the training capabilities, educational institutions of excellence such as IITs, IIMs, etc., can develop and provide tailored course content on AI in finance. Scalable and inclusive capacity-building models and programs must also be developed to reach a wider base of the workforce, particularly in smaller institutions and rural branches.

Recommendation 10 – Capacity Building within REs: REs should develop AI-related capacity and governance competencies for the Board and C suite, as well as structured and continuous training, upskilling, and reskilling programs across the broader workforce who use AI, to effectively mitigate AI risks and guide ethical as well as ensure responsible AI adoption. ***[REs, Medium term]***

4.4.38 Developing Capacity for Financial Sector Regulators and Supervisors:

Regulators and supervisors must also develop an understanding of AI technologies, their innovation potential and the ethical challenges they pose. Without it, regulators may inadvertently curtail innovation, issue policies, or adopt supervisory approaches that either overlook critical challenges or fail to provide appropriate safeguards. This gap could result in ineffective oversight, regulatory blind spots, or missed innovation opportunities. To address this challenge, regulators and supervisors must strengthen their institutional capacity through structured and continuous training focused on the evolving landscape of AI, which is expected to ensure that regulatory responses and supervisory oversight remain relevant and proportionate to the dynamic nature of AI deployment in financial services. RBI may consider establishing an AI institute for the financial sector to support capacity building for regulators and supervisors. The AI institute should also conduct industry-training programmes and research activities on emerging AI trends, thereby enabling more responsible AI adoption across the broader financial ecosystem.

Recommendation 11 – Capacity Building for Regulators and Supervisors:
Regulators and supervisors should invest in training and institutional capacity building initiatives to ensure that they possess an adequate understanding of AI technologies and to ensure that the regulatory and supervisory frameworks match the evolving landscape of AI, including associated risks and ethical considerations. RBI may consider establishing a dedicated AI institute to support sector-wide capacity development.

[RBI, Medium term]

4.4.39 Establishing a Framework for Sharing Best Practices and Lessons on AI

Use Cases and Adoption: Once AI innovation has been successfully catalysed across the length and breadth of the financial sector, it will be important to put in place a structured framework to share experiences and lessons, opportunities to replicate success, avoid common pitfalls, and identify emerging risks. Regular workshops, policy dialogues, and discussions will keep the sector updated on new developments and opportunities for AI adoption. A voluntary and industry-driven framework will make it possible for the sector to learn from each other's experience on what works, what doesn't, and what warrants regulatory scrutiny, while at the same time, positioning India as a global hub for AI-driven financial innovation.

Recommendation 12 – Framework for Sharing Best Practices: *The financial services industry, through bodies such as IBA or SROs, should establish a framework for the exchange of AI-related use cases, lessons learned, and best practices and promote responsible scaling by highlighting positive outcomes, challenges, and sound governance frameworks.*

[Industry Association / SRO, Medium term]

4.4.40 Fostering Responsible Innovation through Recognition or Rewards:

Another way to build capacity is to encourage innovation and experimentation by putting in place carefully designed recognition and incentive frameworks. This could include initiatives such as an annual ‘AI in Finance Award’ to recognise exemplary AI innovations across categories like financial inclusion, customer service, fraud detection, AI compliance toolkits, etc. Regulators and industry bodies could institute periodic ‘AI Challenge Grants’ or ‘AI Innovation Prizes’ to incentivise the development of cutting-edge AI solutions. By fostering competitive innovation, especially among non-regulated entities such as start-ups and smaller firms that may lack visibility and resources, it will encourage these organisations to focus on internal capacity building.

Recommendation 13 – Recognise and Reward Responsible AI Innovation: *Regulators and industry bodies should introduce structured programs to recognise and reward responsible AI innovation in the financial sector, particularly those that demonstrate positive social impact and embed ethical considerations by design.*

[Regulators and Industry, Medium term]

Risk Mitigation Framework

4.4.41 While it is important to enable AI innovation, one cannot lose sight of the risks that could arise as AI starts to get increasingly integrated into the financial sector. To this end, it is just as important to put in place a risk mitigation framework that implements the safeguards necessary for ensuring that AI is deployed in a safe and responsible manner. The following recommendations are designed to ensure that AI risks are managed and mitigated appropriately and are presented across three distinct pillars: Governance, Protection and Assurance.

Governance Pillar

4.4.42 Innovation thrives when it operates within a framework of transparent and accountable governance structure. Governance serves as the backbone of any AI-

related risk management strategy, ensuring that all AI initiatives align with regulatory expectations, ethical principles, and business objectives.

4.4.43 Establish a Board-Approved AI Policy within REs: AI adoption in financial institutions often takes place without a consistent organisational stance on what constitutes responsible or ethical use. In the absence of a formal policy, different teams within the same organisation may proceed with different interpretations as to what constitutes acceptable risk. This could lead to fragmented implementation, blind spots, and consumer harm. It also risks leaving the board and senior management unaware of the risks or reputational consequences of their use of AI.

4.4.44 Just as financial institutions have board-approved policies on credit, cybersecurity, or outsourcing, they should put in place a board-approved AI policy that explicitly articulates the institution's position on AI governance, ethics, and accountability that is aligned with its values, obligations, and risk appetite. The policy should also include a clear risk classification framework for AI use cases, categorizing them as low risk, medium risk, or high risk depending on factors such as impact on customers, criticality, and potential for harm. An indicative classification could be as follows: Low-risk use cases may include internal applications such as document summarisation, email classification, etc., where the outcomes have limited impact. Medium-risk use cases could involve customer-facing tools like chatbots, fraud detection systems, etc., where AI is used for preliminary assistance. High-risk use cases would include critical functions such as credit underwriting, autonomous AI systems that handle customer interactions, make financial decisions, or move customer funds, where errors could have significant consequences for customers or the financial system. Importantly, REs must periodically review and update these classifications to ensure they remain relevant and responsive to the evolving situations.

4.4.45 It should be the responsibility of the Risk Management Committee or similar body to identify, assess, and mitigate AI-related risks and integrate them into the institution's overall risk mitigation framework. Additionally, it could consider putting in place an AI Adoption Committee or leveraging any existing body tasked with technology adoption to bridge functional teams across business, risk, compliance, and technology departments, ensuring that AI innovation and adoption are cross-

departmental and well managed. All functionaries responsible for risk must be well equipped to explicitly incorporate AI-related risks into the organization's risk mitigation framework.

4.4.46 The use of third-party, or off-the-shelf AI tools (e.g., generative AI applications) for official purposes, such as drafting documents, report summarisation, data analysis etc., should be governed by the policies of the organisation. REs must ensure that their internal AI policies are compliant with the broader national AI governance and regulatory frameworks. A draft AI policy template could be prepared by industry bodies/ SROs so that smaller entities that may not have the skillset to develop one from scratch could adapt it to suit their specific organizational needs. A suggested outline of a Board-approved policy on AI has been provided in Annexure V for reference.

Recommendation 14 – Board-Approved AI Policy: *To ensure the safe and responsible adoption of AI within institutions, REs should establish a board-approved AI policy which covers key areas such as governance structure, accountability, risk appetite, operational safeguards, auditability, consumer protection measures, AI disclosures, model life cycle framework, and liability framework. Industry bodies should support smaller entities with an indicative policy template.*

[REs and Industry, Medium term]

4.4.47 Governing the AI Data Lifecycle: High-quality data is key to trustworthy and effective AI systems. However, weak internal controls relating to access, usage, and storage of data could amplify biases, reduce performance, and result in unreliable outcomes. Robust governance processes at the institutional level complement national policies by building operational trust and enabling safe AI deployment. Accordingly, establishing robust internal data governance frameworks across the entire data lifecycle becomes paramount. From the point of data collection to its final deletion or archival, each stage must be governed by clear internal policies. Data used for AI applications must be relevant, fairly representative, and ethically sourced. Weak controls at any stage, whether due to poor quality checks or failure to adhere to consent obligations, can undermine the integrity of AI systems and expose institutions to reputational, legal, and operational risks. REs should put in place guardrails, especially when using open source or external AI models, to ensure that sensitive

customer and institutional data remains within secure environments under the control of the institution. The Digital Personal Data Protection (DPDP) Act provides overarching principles for data protection and privacy and REs are obligated to adhere to DPDP Act provisions and operationalise responsible data management.

Recommendation 15 – Data Lifecycle Governance: *REs must establish robust data governance frameworks, including internal controls and policies for data collection, access, usage, retention, and deletion for AI systems. These frameworks should ensure compliance with the applicable legislations, such as the DPDP Act, throughout the data life cycle.* **[REs, Medium term]**

4.4.48 Establishing an AI System Governance Framework for Safe and Compliant AI Development: AI system governance refers to the structured oversight of AI models and systems, including both conventional AI models and increasingly autonomous AI systems, supported by clear policies, roles, and controls. Robust model governance is critical to ensuring the reliability, safety, and accountability of AI systems. REs must implement appropriate governance mechanisms across the entire AI model lifecycle, covering model design, development, deployment, and decommissioning. REs should maintain a model inventory and documentation that records essential details, including objectives, design features, usage context, performance benchmarks, intended outcomes, etc. Models, whether developed internally or sourced externally, should undergo rigorous validation and periodic testing to ensure they perform as intended. REs must put in place mechanisms to detect and address issues such as model degradation, model drift, bias, or unexplained behaviour, with clearly defined fallback mechanisms. Ongoing performance monitoring, internal audits, and red-teaming exercises should be employed to identify and subsequently rectify vulnerabilities. Any errant model behaviour or incidents must be formally recorded and reported through appropriate channels. REs should also establish procedures for the timely winding down or replacement of models that become outdated or non-compliant.

4.4.49 Emerging developments in AI have given rise to increasingly autonomous systems that allow AI applications to independently execute tasks that would otherwise have required human involvement. When these systems are tasked with financial functions such as investment decisions, loan processing, or payment execution, they

are able to operate with access to real-world customer assets like bank accounts or financial data. While this presents opportunities for efficiency and scale, it also introduces significant risks. Autonomous AI, even when performing simple individual tasks, can generate complex, unintended consequences if not managed well. REs must use autonomous AI only after establishing clear safeguards and accountability frameworks, supported by well-defined testing protocols and standard operating procedures (SoPs). Consumers should be made to fully understand the consequences before being allowed to use such tools. While exceptions may be considered for the use of autonomous AI in routine or low-risk tasks, human oversight remains a critical factor in medium-risk to high-risk tasks. REs must clearly define the tasks AI can perform autonomously and instances when human oversight is required. REs must remain liable for the actions and outcomes of the autonomous AI systems they deploy, just as they are for other forms of operational or technological risk.

Recommendation 16 – AI System Governance Framework: *REs must implement robust model governance mechanisms covering the entire AI model lifecycle, including model design, development, deployment, and decommissioning. Model documentation, validation, and ongoing monitoring, including mechanisms to detect and address model drift and degradation, should be carried out to ensure safe usage. REs should also put in place strong governance before deploying autonomous AI systems that are capable of acting independently in financial decision-making. Given the higher potential for real-world consequences, this should include human oversight, especially for medium and high-risk use cases and applications. [REs, Medium term]*

4.4.50 AI Specific Evaluations in the Product Approval Processes for AI-Enabled Products and Solutions: As AI-enabled products and solutions are increasingly used in financial services, there is a risk that existing product approval mechanisms may be inadequate to identify and address AI-specific risks. There is a need to integrate AI-specific evaluations into these approval processes.

4.4.51 AI-specific risk evaluations should address key elements such as fairness, bias, understandability, customer protection, cybersecurity, and compliance across the entire product lifecycle from pre-development to deployment and use. The product approval process should assess the quality of data, exclusion of sensitive attributes, data pre-processing, random sampling of outputs, back testing, subject matter expert

review, feedback mechanisms, etc. REs are encouraged to deploy internal AI sandboxes to enable controlled testing and validation of their models before deployment. To ensure objectivity, the product approval evaluation team should be independent from the teams responsible for AI model development and deployment.

Recommendation 17 – Product Approval Process: *REs should ensure that all AI-enabled products and solutions are brought within the scope of the institutional product approval framework, and that AI-specific risk evaluations are included in the product approval frameworks.* ***[REs, Medium term]***

Protection Pillar

4.4.52 In an AI-driven financial ecosystem, the protection of data, confidential information and consumer interests is paramount to building trust and resilience. Putting in place these protections will ensure that consumers are not harmed while using AI systems.

4.4.53 Putting Consumers First and Safeguarding the Consumer Experience:

The failure to proactively address AI risks not only harms individuals but also erodes public trust in AI innovations. REs need to establish robust, board-approved consumer protection frameworks that focus on transparency, fairness, and provide clear recourse mechanisms.

4.4.54 Consumers must have effective means of grievance redressal with regard to their interactions with AI or decisions made by AI. REs must embed clear and accessible safeguards into all AI-enabled offerings. Consumers must be explicitly informed whenever they are interacting with AI systems and should always have the option to switch to human representatives when they want. Firms should not be allowed to deceive customers by falsely claiming to be using AI. REs should ensure that AI-driven systems operate only through secure and verifiable channels such as verified 1601 series phone numbers for voice interactions, watermarked digital interfaces for online channels, and clearly labelled platforms.

4.4.55 Consumers should be able to escalate any AI-related issues to human representatives through easily accessible and effective processes. REs must launch targeted and activity-based awareness campaigns that inform customers about their rights when interacting with AI, explain how AI is being utilized in financial services,

and detail the grievance redressal mechanisms that are available to them. Trust is not built by technology alone; it is earned by putting people first.

Recommendation 18 – Consumer Protection: *REs should establish a board-approved consumer protection framework that prioritises transparency, fairness, and accessible recourse mechanisms for customers. REs must invest in ongoing education campaigns to raise consumer awareness regarding safe AI usage and their rights.*

[REs, Medium term]

4.4.56 Mitigating Cybersecurity Threats: The adoption of AI in financial services introduces new cybersecurity risks. AI models potentially expand the attack surface, exposing institutions to threats such as adversarial attacks, data poisoning, and model manipulation. Malicious actors can use AI to automate phishing, create deepfake frauds, and conduct intelligent cyber intrusions. The use of AI by attackers can significantly reduce the time required to conduct cyberattacks and increase their volume.

4.4.57 REs deploying AI in high-risk use cases or using AI extensively in their products and processes should identify vulnerabilities, adversarial weaknesses, and potential security risks before deployment. Cybersecurity assessments must not be limited to the testing or pre-deployment phase, but instead should be a continuous process even after models have been deployed.

4.4.58 AI also offers powerful tools to strengthen cybersecurity. AI-driven anomaly detection, predictive threat intelligence, real-time intrusion monitoring, and adaptive defence systems can significantly enhance the resilience of financial institutions. Additionally, AI systems should be capable of being terminated instantly if there is a risk of significant harm. Consumers should be educated regarding the potential cybersecurity risks involved in the use of AI.

Recommendation 19 – Cybersecurity Measures: *REs must identify potential security risks on account of their use of AI and strengthen their cybersecurity ecosystems (hardware, software, processes) to address them. REs may also make use of AI tools to strengthen cybersecurity, including dynamic threat detection and response mechanisms.*

[REs, Medium term]

4.4.59 Red Teaming of AI Models and Applications: A key challenge in AI deployment is that the harm caused is sometimes only visible after it has affected several people. A proactive way to address this problem is structured red teaming, an adversarial testing approach designed to challenge AI systems to reveal hidden vulnerabilities, stress points, and risks. For instance, investigating if the AI model memorises and inadvertently leaks sensitive data such as account numbers or transaction details when queried in unintended ways.

4.4.60 Since red teaming is proactive, it makes it possible for REs to anticipate failures and mitigate them in advance. This strengthens model resilience, prevents cascading failures, and enhances consumer and system-level trust in AI-enabled financial services. REs, which deploy medium-risk and high-risk AI applications, should make red teaming a regular practice conducted at periodic (at least semi-annual) intervals. For low-risk AI applications, red teaming should at least be conducted at the pre-deployment stage. In addition, red teaming should be carried out before all major model updates, after vulnerabilities have been detected, when there has been a change in the operational environment, or in the event of evolving regulatory requirements. Findings from red teaming exercises should be documented and made accessible to the audit/ supervisory teams, along with steps taken to mitigate them. Key insights should be shared as part of broader knowledge dissemination efforts across the ecosystem to support collective risk awareness and capacity building.

Recommendation 20 – Red Teaming: REs should establish structured red teaming processes that span the entire AI lifecycle. The frequency and intensity of red teaming should be proportionate to the assessed risk level and potential impact of the AI application, with higher risk models being subject to more frequent and comprehensive red teaming. Trigger-based red teaming should also be considered to address evolving threats and changes.

[REs, Medium term]

4.4.61 Ensuring Business Continuity of AI Systems: Despite robust controls, testing, etc., AI systems can fail. Resilience lies in the rapid detection of issues, transparent remediation, and systemic learning. Institutions must embed AI-specific contingencies within their operational resilience frameworks. Failures fall into two categories: traditional system failures (e.g., server outages, cyber incidents) that can be managed via standard Business Continuity Plans (BCPs); and AI-specific failures,

where models remain functional but produce unreliable outputs due to distribution shifts or evolving input-output mappings. For instance, a biased model may continue to deny service to a particular segment. In the case of such AI-specific failures, the challenge often lies in the fact that most models are trained with the assumption that the data encountered during deployment will closely resemble the data used during training. When this assumption fails, the model may continue to run without raising any flags, even as its outputs become increasingly inaccurate.

4.4.62 AI-specific BCPs must go beyond traditional recovery strategies and incorporate fallback mechanisms tailored to AI failure modes. This includes having safeguards and fallback mechanisms such as mandatory human-in-the-loop reviews and continuous model performance monitoring. An AI model should be able to declare itself “unavailable” when it fails and trigger backup processes. Institutions should also conduct regular BCP drills in relation to AI-specific failures, simulating scenarios such as data drift and concept drift and implementing periodic human validation checks on a sample of AI decisions (e.g., 1%) to detect silent model degradation.

Recommendation 21 – Business Continuity Plan for AI Systems: REs must augment their existing BCP frameworks to include both traditional system failures as well as AI model-specific performance degradation. REs should establish fallback mechanisms and periodically test the fallback workflows and AI model resilience through BCP drills. ***[REs, Medium term]***

4.4.63 AI Incident Reporting for REs and Sectoral Risk Intelligence Framework:

AI-related incidents in the financial sector can arise across use cases, often reflecting known failure modes of AI systems, such as bias, lack of explainability, privacy breaches, or unintended actions. For instance, AI models used for credit, loan, or insurance decisions may exhibit bias against specific demographic groups; fraud detection models may be circumvented by novel attack strategies; AI agents may act beyond their intended scope; Chatbots may misinterpret customer inputs and conflict with ethical principles. Such incidents can result in significant financial, operational, or reputational harm. Without structured reporting and analysis, such risks may remain hidden until they cause systemic harm. A tiered incident reporting framework is essential at the entity, sector, and national levels to identify patterns, address vulnerabilities, and prevent recurrence. Inspired by the aviation industry, the financial

sector must promptly report incidents as soon as possible. The time within which reporting needs to be done could vary based on severity and system-wide implications.

4.4.64 Regulators should design a system for aggregation of risk data for macro-level insights, possibly via an expanded Emerging Technology (EmTech) Repository, and encourage transparent, non-punitive reporting. At the national level, analysis of such incident data should be channelled into inter-regulatory coordination forums to inform coordinated responses and strengthen sectoral resilience. Reporting what was observed, where and how it went wrong, and what remedial action was taken should be sufficient to enable shared learning. This will also serve as an early warning system and ensure AI adoption remains resilient, inclusive, and grounded in public trust. Where a customer has been adversely impacted, compensation must be provided by the RE, but reporting such incidents should not, by itself, trigger penal action if timely corrective measures have been taken and disclosure is complete. The objective is to foster a culture of early AI incident reporting across the financial sector, so that entities and the broader sector can adapt accordingly. An indicative sample Incident Reporting Form template has been provided in Annexure VI for reference.

Recommendation 22 – AI Incident Reporting and Sectoral Risk Intelligence Framework: *Financial sector regulators should establish a dedicated AI incident reporting framework for REs and FinTechs and encourage timely detection and reporting of AI-related incidents. The framework should adopt a tolerant, good-faith approach to encourage timely disclosure.* **[REs, Regulators Medium term]**

Assurance Pillar

4.4.65 The assurance pillar is designed to provide oversight throughout the AI lifecycle. It focuses on monitoring emerging risks and feeding those insights back into both institutional and system-wide responses. The framework addresses key questions such as:

- Do organisations have visibility over the AI systems in use, and do they provide transparent disclosures to stakeholders? Do policymakers have visibility over sector-wide AI adoption and the potential risks building up?

- Are there effective controls through proactive and continuous testing and auditing of the AI systems to ensure they behave in line with established principles and guidelines?
- Is there a clear and fair liability framework that ensures accountability for errors and failures, while also ensuring innovation is not stifled?

4.4.66 Trust in AI cannot be assumed; it must be built and, more importantly, sustained. The assurance pillar is the mechanism through which trust is continually reinforced and upheld.

4.4.67 Creating Visibility – Maintaining an AI Inventory within Institutions and a Sector-Wide AI Repository: A key challenge in assessing AI risks is the lack of clear visibility as to where and how AI systems are deployed within institutions and across the sector. Without a structured and updated view of AI usage, it becomes difficult for institutions and supervisors to assess risk exposure or monitor changes over time.

4.4.68 REs should maintain a comprehensive inventory of AI systems in use across their operations. Among other aspects, the inventory should include information on:

- **AI Models and Algorithms:** All AI models in use, including the type of model (e.g., machine learning, deep learning, natural language processing, GenAI), their purpose, and the functional areas they support.
- **Use Cases and Applications:** A clear description of how each model is deployed.
- **Dependencies:** Third-party providers, cloud service providers, data sources, and any other external components that can influence AI model performance.
- **Risk Categorisation:** Assessment of each AI system's risk level (e.g. High, Medium, Low) based on the board-approved policies.
- **Grievances:** A record of the volume and nature of grievances filed in respect of these AI systems and how they were resolved. This includes information as to whether the AI solutions have been modified in response to user complaints.

4.4.69 This inventory should be updated semi-annually and must be readily available for supervisory inspections, audits, and ongoing risk monitoring efforts. Maintaining this AI inventory will give both the REs and supervisors a view of where and how AI/ML models are being used to better categorise risk, improve oversight, and ensure responsible deployment.

4.4.70 To complement this institution-level visibility, a sector-wide AI repository should also be developed to collect and maintain aggregate information on AI models and applications across all REs. The EmTech Repository of RBIH can be leveraged by expanding its scope for this purpose. The repository should capture bare minimum, indicative information such as the types and number of AI models deployed across institutions, high-risk use cases, critical dependencies, incidents of AI model failures, ethical breaches, etc. This can also help in monitoring systemic risks such as model correlation and model herding, which, if left unchecked, can amplify vulnerabilities across institutions and potentially pose broader financial stability risks.

4.4.71 Over time, associations such as IBA or SROs can consider developing a Responsible AI Adoption Score or Index for the financial sector, which could serve as a baseline measure to track the maturity and ethical integration of AI across the sector.

Recommendation 23 – AI Inventory within REs and Sector-Wide Repository: REs should maintain a comprehensive, internal AI inventory that includes all models, use cases, target groups, dependencies, risks and grievances, updated at least half yearly, and it must be made available for supervisory inspections and audits. In parallel, regulators should establish a sector-wide AI repository that tracks AI adoption trends, concentration risks, and systemic vulnerabilities across the financial system with due anonymisation of entity details. ***[Regulators and REs, Short term]***

4.4.72 Ensuring Responsible AI through a comprehensive Audit Framework: Audits help to independently confirm that systems are operating as intended and within regulatory boundaries. Unlike traditional systems, AI systems are often non-deterministic, adaptive, and opaque, making it difficult to evaluate whether the output is consistent, fair, and compliant with internal policies. This also makes AI prone to specific risks, such as biases and data drift. In order to cater to this, the audit framework that is put in place needs to be risk-based and proportionate in order to ensure that AI systems operate within the guardrails set by the regulator while still allowing room for innovation.

4.4.73 The audit should aim to verify not only that the system works technically but also that it aligns with the 7 *Sutras*. An effective AI audit should cover:

- **Input Data Audit:** It needs to certify that the data used for training or inference is accurate, unbiased, and collected in conformity with the data regulations.

- **Model and Algorithm Audit:** It needs to certify that the model architecture, training methods, and decision logic align with the intended purpose and that the models are resilient against manipulation or misuse that could cause them to act contrary to their stated objectives.
- **Output and Behaviour Audit:** It needs to certify that the decisions made by the AI model, such as approving a loan, flagging a transaction, or responding to a customer, are explainable, fair, consistent, and compliant with the applicable guidelines and principles and that there are safeguards in place to ensure these outputs cannot be misused or manipulated by bad actors.

4.4.74 The audit should be tailored to the risk level of each application. For instance, internal AI audits for low-risk use cases (e.g., document summarisation) may be minimal, while audits of high-risk applications (e.g., such as credit decisioning) should be detailed. The audit should also confirm that mechanisms exist to stop, pause or unwind AI-driven processes in a controlled manner in case of malfunction or policy breach. It should also verify the presence of a Business Continuity Plan (BCP) for core AI systems and ensure that human oversight and override mechanisms are available for critical decisions. In cases where the risk is particularly high or where internal expertise may be limited, third-party AI audits by independent experts can provide necessary assurance. Audits should be periodic and evolving, with mechanisms to continuously update audit controls and coverage areas, considering new risks, such as agent-to-agent interactions.

4.4.75 Supervisory audits should also evolve accordingly. Inspection by supervisors should include standardised AI-specific checklists and model risk templates tailored to AI systems, providing clarity on what aspects to audit, how to evaluate performance, and how institutions can demonstrate compliance.

Recommendation 24 - AI Audit Framework: REs should implement a comprehensive, risk-based, calibrated AI audit framework, aligned with a board-approved AI risk categorisation, to ensure responsible adoption across the AI lifecycle, covering data inputs, model and algorithm, and the decision outputs.

a. Internal Audits: As the first level, REs should conduct internal audits proportionate to the risk level of AI applications.

b. Third-Party Audits: *For high-risk or complex AI use cases, independent third-party audits should be undertaken.*

c. Periodic Review: *The overall audit framework should be reviewed and updated at least biennially to incorporate emerging risks, technologies, and regulatory developments.*

Supervisors should also develop AI-specific audit frameworks, with clear guidance on what to audit, how to assess it, and how to demonstrate compliance.

[Supervisors and REs, Medium term]

4.4.76 Promoting Transparency through Public Disclosures of AI Use and Safeguards:

In order to foster public trust and provide assurance, customers and external stakeholders should have visibility into how AI is being governed and whether their concerns are being acknowledged and addressed. To this end, having AI disclosures in publicly available reports can play a vital role in strengthening confidence among the public and stakeholders. Not only will this help to promote market discipline, it will also nudge institutions towards responsible AI practices. Just as climate risk and cybersecurity disclosures are now part of annual reports and ESG filings, AI disclosures should become a regular feature of REs' annual reports. The disclosures may contain necessary details regarding AI governance frameworks in place, adoption areas, ethical guidelines adopted, consumer protection measures, complaints and grievances handled, etc.

Recommendation 25 – Disclosures by REs: *REs should include AI-related disclosures in their annual reports and websites. Regulators should specify an AI-specific disclosure framework to ensure consistency and adequacy of information across institutions.*

[REs, Regulators, Short term]

4.4.77 Enabling Responsible AI Compliance through Standardised Assessment

Toolkits: REs may lack standardised, practical mechanisms to demonstrate that their AI systems are performing in line with the 7 Sutras. By making available standardised open-source tools which can evaluate the AI model from different dimensions, such as model accuracy, transparency, fairness, etc., REs will be able to demonstrate compliance.

4.4.78 Regulators should facilitate the development of industry-led AI Compliance Toolkits to help REs validate that their AI models and applications meet regulatory expectations. These toolkits can serve both as a diagnostic as well as a benchmarking mechanism, enabling the validation of key AI risks. The use of the toolkit should be voluntary but strongly encouraged, especially for smaller and mid-sized REs that may lack internal capabilities. The toolkit could be developed and maintained by an industry body or SRO, or a consortium of financial sector participants. Third-party service providers should be encouraged to offer toolkit-based validation services, without regulatory endorsement. Regulators may periodically identify and share best practices to guide the continuous improvement of these toolkits. The toolkits would offer a baseline confidence level but not absolve institutions of their responsibility for end-to-end AI risk management, and are intended to complement, not replace, internal validations or oversight.

Recommendation 26 – AI Toolkit: *AI Compliance Toolkit will help REs validate, benchmark, and demonstrate compliance against key responsible AI principles such as fairness, transparency, accountability, and robustness. The toolkit should be developed and maintained by a recognised SRO or industry body.*

[Regulators and Industry, Medium term]

4.5 Conclusion - Weaving It All Together

4.5.1 As AI continues to evolve and reshape the financial landscape, it brings with it both transformative opportunities and complex challenges. This report has sought to present a balanced and forward-looking framework of how AI can be responsibly and ethically enabled in the Indian financial sector. At the heart of the FREE-AI framework are the 7 *Sutras*, the foundational principles which are the living spirit of the framework. The 6 Pillars provide structural balance by enabling innovation as well as mitigating risks. Finally, the 26 Recommendations bring it all to life with specific, implementable steps that translate aspiration into action. The recommendations have been carefully crafted to embody and advance the *Sutras*. Together, the *Sutras*, the Pillars, and the Recommendations, forge a progressive path forward for all stakeholders, including regulators, financial institutions, technology service providers, to harness the potential of AI in the financial sector.

Summary of Sutras and Recommendations

Summary of the 7 Sutras	
Sl. No.	Description
1	Trust is the Foundation: Trust is non-negotiable and should remain uncompromised
2	People First: AI should augment human decision-making but defer to human judgment and citizen interest
3	Innovation over Restraint: Foster responsible innovation with purpose
4	Fairness and Equity: AI outcomes should be fair and non-discriminatory
5	Accountability: Accountability rests with the entities deploying AI
6	Understandable by Design: Ensure explainability for trust
7	Safety, Resilience, and Sustainability: AI systems should be secure, resilient and energy efficient

Summary of Recommendations		
Sl. No.	Description	Action and Timeline
Innovation Enablement Framework		
	Infrastructure Pillar	
1	Financial Sector Data Infrastructure: A high-quality financial sector data infrastructure should be established, as a digital public infrastructure, to help build trustworthy AI models for the financial sector. It may be integrated with the AI Kosh – India Datasets Platform, established under the IndiaAI Mission.	Regulators and Government, Short term
2	AI Innovation Sandbox: An AI innovation sandbox for the financial sector should be established to enable REs, FinTechs, and other innovators to develop AI-driven solutions, algorithms, and models in a secure and controlled environment. Other FSRs should also collaborate to contribute to and benefit from this initiative.	Regulators RBI, MeitY, FSRs, Short term

3	Incentives and Funding Support: Appropriate incentive structures and infrastructure must be put in place to encourage inclusive and equitable AI usage among smaller entities. To support innovation and to meet strategic sectoral needs, RBI may also consider allocating a fund for setting up of data, compute infrastructure.	RBI and Government, Medium term
4	Indigenous Financial Sector Specific AI Models: Indigenous AI models (including LLMs, SLMs, or non LLM models) tailored specifically for the financial sector should be developed and offered as a public good.	Regulators, SROs and Industry, Medium term
5	Integrating AI with DPI: An enabling framework should be established to integrate AI with DPI in order to accelerate the delivery of inclusive, affordable financial services at scale.	Regulators, Medium term
	Policy Pillar	
6	Adaptive and Enabling Policies: Regulators should periodically undertake an assessment of existing policies and legal frameworks to ensure they effectively enable the AI-driven innovations and address the AI-specific risks. Regulators should develop a comprehensive AI policy framework for the financial sector, anchored in the Committee's 7 <i>Sutras</i> to provide flexible, forward-looking guidance for AI innovation, adoption, and risk mitigation across the sector. The RBI may consider issuing consolidated AI Guidance to serve as a single point of reference for regulated entities and the broader FinTech ecosystem on the responsible design, development, and deployment of AI solutions.	RBI, Medium term
7	Enabling AI-Based Affirmative Action: Regulators should encourage AI-driven innovation that accelerates financial inclusion of underserved and unserved sections of society and other such affirmative actions by lowering compliance expectations as far as is possible, without compromising basic safeguards.	Regulators, Medium term

8	<p>AI Liability Framework: Since AI systems are probabilistic and non-deterministic, regulators should adopt a graded liability framework that encourages responsible innovation. While REs must continue to remain liable for any loss suffered by customers, an accommodative supervisory approach where the RE has followed appropriate safety mechanisms such as incident reporting, audits, red teaming etc., is recommended. This tolerant supervisory stance should be limited to first time / one-off aberrations and denied in the event of repeated breaches, gross negligence, or failure to remediate identified issues.</p>	Regulators, Medium term
9	<p>AI Institutional Framework: A permanent multi-stakeholder AI Standing Committee should be constituted under the Reserve Bank of India to continuously advise it on emerging opportunities and risks, monitor the evolution of AI technology, and assess the ongoing relevance of current regulatory frameworks. The Committee may be constituted for an initial period of five years, with a built-in review mechanism and a sunset clause. A dedicated institution should be established for the financial sector, operating under a hub-and-spoke model to the national-level AI Safety Institute, for continuous monitoring and sectoral coordination.</p>	Regulators, RBI, Short term
	Capacity Pillar	
10	<p>Capacity Building within REs: REs should develop AI-related capacity and governance competencies for the Board and C suite, as well as structured and continuous training, upskilling, and reskilling programs across the broader workforce who use AI, to effectively mitigate AI risks and guide ethical as well as ensure responsible AI adoption.</p>	REs, Medium term
11	<p>Capacity Building for Regulators and Supervisors: Regulators and supervisors should invest in training and institutional capacity building initiatives to ensure that they possess an adequate understanding of AI technologies and to</p>	RBI, Medium term

	ensure that the regulatory and supervisory frameworks match the evolving landscape of AI, including associated risks and ethical considerations. RBI may consider establishing a dedicated AI institute to support sector-wide capacity development.	
12	Framework for Sharing Best Practices: The financial services industry, through bodies such as IBA or SROs, should establish a framework for the exchange of AI-related use cases, lessons learned, and best practices and promote responsible scaling by highlighting positive outcomes, challenges, and sound governance frameworks.	Industry Association / SRO, Medium term
13	Recognise and Reward Responsible AI Innovation: Regulators and industry bodies should introduce structured programs to recognise and reward responsible AI innovation in the financial sector, particularly those that demonstrate positive social impact and embed ethical considerations by design.	Regulators and Industry, Medium term
Risk Mitigation Framework		
	Governance Pillar	
14	Board Approved AI Policy: To ensure the safe and responsible adoption of AI within institutions, REs should establish a board-approved AI policy which covers key areas such as governance structure, accountability, risk appetite, operational safeguards, auditability, consumer protection measures, AI disclosures, model life cycle framework, and liability framework. Industry bodies should support smaller entities with an indicative policy template.	REs and Industry, Medium term
15	Data Lifecycle Governance: REs must establish robust data governance frameworks, including internal controls and policies for data collection, access, usage, retention, and deletion for AI systems. These frameworks should ensure compliance with the applicable legislations, such as the DPDP Act, throughout the data life cycle.	REs, Medium term

16	AI System Governance Framework: REs must implement robust model governance mechanisms covering the entire AI model lifecycle, including model design, development, deployment, and decommissioning. Model documentation, validation, and ongoing monitoring, including mechanisms to detect and address model drift and degradation, should be carried out to ensure safe usage. REs should also put in place strong governance before deploying autonomous AI systems that are capable of acting independently in financial decision-making. Given the higher potential for real world consequences, this should include human oversight, especially for medium and high-risk use cases and applications.	REs, Medium term
17	Product Approval Process: REs should ensure that all AI-enabled products and solutions are brought within the scope of the institutional product approval framework, and that AI-specific risk evaluations are included in the product approval frameworks.	REs, Medium term
	Protection Pillar	
18	Consumer Protection: REs should establish a board-approved consumer protection framework that prioritises transparency, fairness, and accessible recourse mechanisms for customers. REs must invest in ongoing education campaigns to raise consumer awareness regarding safe AI usage and their rights.	REs, Medium term
19	Cybersecurity Measures: REs must identify potential security risks on account of their use of AI and strengthen their cybersecurity ecosystems (hardware, software, processes) to address them. REs may also make use of AI tools to strengthen cybersecurity, including dynamic threat detection and response mechanisms.	REs, Medium term
20	Red Teaming: REs should establish structured red teaming processes that span the entire AI lifecycle. The frequency and	REs, Medium term

	intensity of red teaming should be proportionate to the assessed risk level and potential impact of the AI application, with higher risk models being subject to more frequent and comprehensive red teaming. Trigger-based red teaming should also be considered to address evolving threats and changes.	
21	Business Continuity Plan for AI Systems: REs must augment their existing BCP frameworks to include both traditional system failures as well as AI model-specific performance degradation. REs should establish fallback mechanisms and periodically test the fallback workflows and AI model resilience through BCP drills.	REs, Medium term
22	AI Incident Reporting and Sectoral Risk Intelligence Framework: Financial sector regulators should establish a dedicated AI incident reporting framework for REs and FinTechs and encourage timely detection and reporting of AI-related incidents. The framework should adopt a tolerant, good-faith approach to encourage timely disclosure.	REs, Regulators Medium term
	Assurance Pillar	
23	AI Inventory within REs and Sector-Wide Repository: REs should maintain a comprehensive, internal AI inventory that includes all models, use cases, target groups, dependencies, risks and grievances, updated at least half yearly, and it must be made available for supervisory inspections and audits. In parallel, regulators should establish a sector-wide AI repository that tracks AI adoption trends, concentration risks, and systemic vulnerabilities across the financial system with due anonymisation of entity details.	Regulators and REs, Short term
24	AI Audit Framework: REs should implement a comprehensive, risk-based, calibrated AI audit framework, aligned with a board-approved AI risk categorisation, to ensure responsible adoption across the AI lifecycle, covering data inputs, model and algorithm, and the decision outputs.	Supervisors and REs, Medium term

	<p>a. Internal Audits: As the first level, REs should conduct internal audits proportionate to the risk level of AI applications.</p> <p>b. Third-Party Audits: For high risk or complex AI use cases, independent third-party audits should be undertaken.</p> <p>c. Periodic Review: The overall audit framework should be reviewed and updated at least biennially to incorporate emerging risks, technologies, and regulatory developments. Supervisors should also develop AI-specific audit frameworks, with clear guidance on what to audit, how to assess it, and how to demonstrate compliance.</p>	
25	<p>Disclosures by REs: REs should include AI-related disclosures in their annual reports and websites. Regulators should specify an AI-specific disclosure framework to ensure consistency and adequacy of information across institutions.</p>	REs, Regulators, Short term
26	<p>AI Toolkit: AI Compliance Toolkit will help REs validate, benchmark, and demonstrate compliance against key responsible AI principles such as fairness, transparency, accountability, and robustness. The toolkit should be developed and maintained by a recognised SRO or industry body.</p>	Regulators and Industry, Medium term

Annexure I – Interactions with Stakeholders by the Committee

Sl. No.	Stakeholder	Participating Individuals and Designation	Date
1	Department of Regulation, RBI	Smt. Usha Janakiraman, Chief General Manager-in-Charge, and team	February 05, 2025
2	Department of Supervision, RBI	Shri Tarun Kumar Singh, Chief General Manager, and team	February 05, 2025
3	National Payment Council of India (NPCI)	Shri Dilip Asbe, Managing Director and CEO, and team	February 05, 2025
4	Black Dot Public Policy Advisors	Shri Mandar Kagade, Founder	February 12, 2025
5	Association of Chartered Certified Accountants (ACCA)	Shri Narayanan Vaidyanathan, Head of Policy Development and Shri Sundeep Jakhar, Head of Public Affairs (India)	February 12, 2025
6	Sarvam AI	Shri Pratyush Kumar, Co-Founder	February 18, 2025
7	Dell Technologies Inc.	Shri Vivek Mohindra, Chief Strategy Officer, Retd. Col Ali Akhtar Jafri, Director Government Affairs and Shri Tabrez Ahmad, Group Director Government Affairs and Public Policy for Asia Pacific and Japan	February 27, 2025
8	Boston Consulting Group (BCG)	Shri Yashraj Erande, Partner and Director, Shri Vikram Khanna, Partner, and members of the global team	March 04, 2025
9	Khaitan & Co	Ms. Vidushi Gupta, Partner, and Ms. Tanu Banerjee, Partner	March 04, 2025
10	Data Security Council of India (DSCI)	Shri Vinayak Godse, Chief Executive Officer	March 19, 2025
11	L&T Finance Limited	Shri Debarag Banerjee, Chief AI and Data Officer	March 20, 2025

12	Vidhi Centre for Legal Policy	Ms. Shehnaz Ahmed, Lead, Applied Law and Technology and Ms Vrinda Pareek, Senior Resident Fellow	March 20, 2025
13	Shri Kris Gopalakrishnan	Chairperson of Axilor Ventures, Chairperson of Reserve Bank of India Innovation Hub, Co-founder, Infosys	March 26, 2025
14	Microsoft Corporation	Ms. Garima Rathore, Tech Policy Lead, and team	April 02, 2025
15	Shri Nandan Nilekani	Co-founder and Chairman of the Board, Infosys, Founding Chairman of the Unique Identification Authority of India	April 04, 2025
16	State Bank of India	Shri Nitin Chugh, Deputy Managing Director and Head of Digital Banking and Transformation, and team	April 5, 2025
17	Accenture Plc	Shri Jayant Prabhu, Managing Director - Data and AI, and team	April 5, 2025
18	Infosys Limited	Shri Ashish Tewari, Head - Responsible AI Office	April 09, 2025
19	Shri Sandeep K Shukla	Professor, Department of Computer Science, Indian Institute of Technology, Kanpur	May 06, 2025

Annexure II – Interactions with Stakeholders by the Secretariat

Sl. No.	Stakeholder	Sl. No.	Stakeholder
1	AU Small Finance Bank Limited	30	Kotak Mahindra Bank Limited
2	Axis Bank Limited	31	Lendingkart Finance
3	Bajaj Finance Limited	32	Mahindra & Mahindra Financial Services Limited
4	Bandhan Bank Limited	33	Navi Finserv Limited
5	Bank of Baroda	34	Neokred
6	Bank of India	35	One Mobikwik Systems Limited
7	Bank of Maharashtra	36	OpenAI Inc
8	BankBazaar	37	Paisabazaar Marketing and consulting private limited
9	Canara Bank	38	Perfios Software Solutions Private Limited
10	CapFloat Financial Services Private Limited	39	PNB Housing Finance Limited
11	Central Bank of India	40	Punjab & Sind Bank
12	Cholamandalam Investment and Finance Company Limited	41	Punjab National Bank
13	Dreamplug Technologies Private Limited (CRED)	42	Razorpay technologies Limited
14	Easebuzz Private Limited	43	Sammaan Capital Ltd
15	Epifi Technologies Private Limited	44	Samunnati Finance Private Limited
16	Ernst & Young LLP	45	Shriram Finance Limited
17	Federal Bank Limited	46	Signzy Technologies Pvt Limited
18	HDB Financial Services Limited	47	Slice Small Finance Bank Limited
19	Hero Fincorp Limited	48	SMFG India Credit company Limited
20	HSBC Limited	49	South Indian Bank Limited
21	ICICI Bank Limited	50	Standard Chartered Bank
22	IDBI Bank Limited	51	State Bank of India
23	IDFC FIRST Bank Limited	52	Tata Consultancy Services Limited
24	Indian Bank	53	UCO Bank
25	Indian Overseas Bank	54	Uni Cards
26	IndusInd Bank Limited	55	Union Bank of India
27	JPMorgan Chase Bank NA	56	YES Bank Limited
28	Karnataka Bank Limited	57	Yubi
29	Karur Vysya Bank Limited	58	Zerodha Capital Private limited

Annexure III – IndiaAI Mission: Strategy and Status

The **IndiaAI Mission** is the Government of India's flagship program to build a cohesive, strategic, and robust AI ecosystem.

The **IndiaAI Compute** pillar focuses on creating a high-end, scalable AI computing ecosystem to deliver Compute-as-a-Service for India's rapidly growing AI startups and research community. So far, over 34,000 GPUs have been made available at subsidized rates through the IndiaAI Compute portal, with an additional 4,000+ GPUs expected in the next phase of empanelment. The mission also plans to establish a government-controlled GPU cluster of about 3,000 GPUs to meet sovereign and strategic needs.

The **IndiaAI Application Development Initiative (IADI)** is designed to foster the development and adoption of at least 25 impactful AI solutions that can drive large-scale socio-economic transformation. The first Innovation Challenge, launched in 2024, wherein thirty applications have advanced to the prototyping phase, with a second round of the challenge set to launch in collaboration with the Ministry of Education.

AIKosh, the IndiaAI Datasets Platform, is envisioned as a unified data platform integrating datasets from government and non-government sources. Launched in beta in March 2025, it currently features over 874 datasets, 207 AI models, and more than 13 development toolkits. The platform has attracted over 265,000 visits, 6,000 registered users, and 13,000+ resource downloads. AIKosh prioritizes data quality scoring, robust search and filtering, Jupyter notebooks for analytics, and secure, permission-based access for contributors.

The **IndiaAI Foundation Models** pillar underscores the importance of building India's own large language models (LLMs) trained on Indian datasets and languages, to ensure sovereign capability and global competitiveness in generative AI. A funding model combining grants and equity support has been introduced, offering 40% of compute costs as grants and taking 60% as equity (via convertible debentures). From 506 proposals received, four startups (Sarvam AI, Soket AI, Gnani AI, and Gan AI) have been selected in the first phase to develop India's foundation models.

The **IndiaAI FutureSkills** pillar is a cornerstone of the mission's human capital strategy, aiming to democratize AI education and build a robust talent pipeline across the country. The program will support 500 PhD fellows, 5,000 Master's students, and 8,000 undergraduates through targeted funding. Research fellowships for PhD scholars are aligned with the Prime Minister's Research Fellowship, offering support of up to ₹55 lakh per fellow. Over 200 students have received fellowships in the first year, with 26 partner institutes onboarding PhD students. Additionally, more than 570 AI and Data Labs are planned nationwide, with 27 labs already in progress and further approvals granted for ITIs and polytechnics across 27 states and UTs.

The **IndiaAI Startup Financing** pillar addresses the critical need for risk capital across the entire lifecycle of AI startups, from prototyping to commercialization. This includes the IndiaAI Startups Global program, launched in collaboration with Station F (Paris) and HEC Paris, which aims to support 10 Indian AI startups in expanding into the European market. A call for proposals to establish state-level Centers of Excellence in AI has also received 29 submissions from 21 states and union territories.

Finally, the **Safe and Trusted AI** pillar seeks to balance innovation with strong governance frameworks to ensure responsible AI adoption. Recognizing India's diverse social, cultural, economic, and linguistic landscape, this pillar focuses on developing contextualized instruments of AI governance. The first Expression of Interest (EoI) selected eight projects addressing themes such as machine unlearning, bias mitigation, privacy-preserving machine learning, explainability, auditing tools, and governance testing frameworks. A second EoI round, focused on watermarking, ethical AI frameworks, risk assessment, stress testing tools, and deepfake detection, received 400+ applications. Plans are also underway to operationalise the IndiaAI Safety Institute under a hub-and-spoke model to address AI risks and safety challenges in collaboration with research institutions and industry partners.

In addition, India is set to host the **AI Impact Summit in February 2026**, building on its role as co-chair of the AI Action Summit and continuing its leadership in shaping global AI discussions.

Annexure IV – AI Specific Enhancements in RBI Master Directions

Existing Regulation	AI Risks Implicitly Covered under	Suggestions for AI Specific Enhancements
1. RBI Guidelines on Outsourcing of Financial Services	(i) Accountability for outsourced activities (ii) Risk management and governance of third-party services	(i) Specific clauses addressing AI specific risks, including algorithmic bias, may be incorporated as applicable into the Outsourcing Agreement. (ii) Specific clauses setting out an obligation to disclose the use of AI by third-party vendors and their subcontractors may be incorporated, as applicable, into the Outsourcing Agreement.
2. Cyber Security Framework in Banks (2016)	(i) Data confidentiality, integrity, and availability (ii) Incident reporting and response mechanisms	(i) May include AI specific threats such as model poisoning and adversarial attacks in the risk assessments under Cyber Security Policy at para 5 of the framework. (ii) May establish protocols for monitoring and mitigating AI related cybersecurity incidents under para 14 of the said framework.
3. Guidelines on Digital Lending dated September 2, 2022	(i) Data privacy and consent in digital lending (ii) Accountability of REs for third-party digital lending apps	(i) May include providing transparency in AI driven credit assessments, including disclosure of use of AI under para 5 “Disclosures to borrowers” of the guidelines. (ii) May include implementation of fairness audits to detect and mitigate algorithmic biases under para 9 “Due diligence requirements with respect to LSPs” of the guidelines.
4. Master Circular on Customer Service in Banks (2015)	(i) Customer rights and grievance redressal mechanisms (ii) Board level oversight of customer service	(i) May include awareness to customer when interacting with AI systems under para 8 “Guidance to customers and Disclosure of Information” of the master circular. (ii) May include establishment of processes for customers to contest

		AI driven decisions under para 16 – “Dealing with Complaints and Improving Customer Relations”
5. Master Direction on Fraud Risk Management (2024)	(i) Framework for early warning signals and fraud detection (ii) Risk management policies approved by the Board	(i) May encourage AI driven fraud detection mechanisms under Chapter III “Framework for Early Warning Signals for Detection of Frauds” (ii) May suggest to regularly test AI models for accuracy and bias in fraud detection under the same Chapter.
6. Master Direction on Information Technology Governance, Risk, Controls and Assurance Practices (2023)	(i) Broad IT governance (ii) Oversight of information systems and related risks	(i) May include Access Control measures for the autonomous AI under para 19 (Access Control) of the Direction.
7. Master Direction on Outsourcing of Information Technology Services (2023)	(i) Risk assessments and due diligence of IT service providers. (ii) Data protection and incident reporting obligations	(i) May be required for the service providers to disclose use of AI in service delivery under para 16 Chapter V (Aspects to be considered in outsourcing agreement) (ii) May amend to include AI specific risk assessments under Chapter IV (Evaluation and Engagement of Service Providers)

Annexure V – Suggested Outline of Board Policy on AI

This document outlines the aspects that an entity should cover while formulating its Board policy on AI. It may be customised to the organisation's needs and complexity of use and aligned with the recommendations of the FREE-AI committee report.

Sl. No.	Section	What needs to be covered
1	Purpose and Scope	<ul style="list-style-type: none"> Define the role of AI in the organisation Identify stakeholders (including, internal departments, external vendors, etc.) State the desired outcomes of AI usage by the organisation
2	Principles	Set out the principles in alignment with the FREE-AI <i>Sutras</i> after taking into account the organisation's needs.
3	Governance Structure and Roles	Define the AI governance structure to address: <ul style="list-style-type: none"> decision-making and accountability, oversight and escalation grievance redressal mechanism
4	AI Life cycle management	<ul style="list-style-type: none"> Ensure rigorous testing and internal product approval processes before deploying any AI application. Conduct continuous red-teaming exercises throughout the AI lifecycle to identify and mitigate emerging risks. List key controls and responsibilities to ensure AI systems are developed and managed in a safe, ethical, and accountable manner through all stages of the AI lifecycle.
5	Data Governance	<ul style="list-style-type: none"> Define the data governance framework for sourcing, cleaning, anonymization, encryption, sharing and purging of data.

	and Management	<ul style="list-style-type: none"> • Address bias detection, synthetic data generation and validation to ensure data integrity and responsible use. • Adhere to open data standards and meta data standards.
6	Risk Management and Controls	<p>Institute processes to identify and mitigate AI risks.</p> <p>Consider the following aspects before deployment:</p> <ul style="list-style-type: none"> • Are the decisions and outputs of the AI application fair, inclusive, accurate and transparent? • Does the AI application introduce new cybersecurity risks or exacerbate the existing cybersecurity risks? • Do customers have adequate recourse for all decisions and outputs of the AI application that could adversely affect their interests? • Does the application adequately convey customers that they are interacting with an AI system? • Does the AI application share or reveal confidential/personal/sensitive information in its interactions
7	Third-Party and Vendor Management	Define the scope, responsibility, and liability of third-party vendors.
8	Policy Review	The AI policy may be reviewed at least once annually to reflect new regulatory requirements and technological developments.

Annexure VI – AI Incident Reporting Form (Indicative Sample)

1. General Information

Date and Time of Incident	
Date and Time of Detection	
Submitting Officer and Contact Details	
Incident Reference ID	

2. Incident Summary

Use Case Involved	<input type="checkbox"/> Credit Decisioning <input type="checkbox"/> Fraud <input type="checkbox"/> Customer Support <input type="checkbox"/> Marketing <input type="checkbox"/> Others: _____
Model Used	<i>e.g, GPT, Llama, or any other internal / proprietary model, LLM/SLM etc.</i>
Is Third-Party Vendor Involved?	<input type="checkbox"/> Yes <input type="checkbox"/> No <i>(If Yes, provide details of vendor)</i>
Brief Description of Incident	<i>(What was the model expected to do? What happened instead? How was it detected?)</i>
Affected Stakeholders	<input type="checkbox"/> Internal <input type="checkbox"/> External <input type="checkbox"/> Both
Estimated Impact	<input type="checkbox"/> Low <input type="checkbox"/> Moderate <input type="checkbox"/> High

3. Additional information

Preliminary Root Cause	<i>Briefly describe suspected cause. Attach detailed report if available</i>
Immediate Response Actions taken	<i>Describe immediate response: model disabled, alerts raised, affected users notified, etc.</i>
Current Status Update	<input type="checkbox"/> Ongoing <input type="checkbox"/> Resolved

References

- 1) Bank of England (BoE) and Financial Conduct Authority (FCA) (2022), *Discussion Paper on Artificial Intelligence and Machine Learning in UK Financial Services*
- 2) Center for Research on Foundation Models (CRFM), (2021), *On the opportunities and risks of foundation models*, Stanford Institute for Human Centered Artificial Intelligence, Stanford University
- 3) Department of Economic Affairs. (2024). *Report of India's G20 Task Force on Digital Public Infrastructure*. Ministry of Finance, Government of India.
- 4) Ernst & Young (2025), *How much productivity can GenAI unlock in India? The Aldea of India: 2025*
- 5) FSB (2017), *Artificial Intelligence and Machine Learning in Financial Services*. 1 November.
- 6) FSB (2024), *The Financial Stability Implications of Artificial Intelligence*, 14 November.
- 7) FCA, (2025), *FCA allows firms to experiment with AI alongside NVIDIA*. Press Release, 9 June
- 8) Frazier, K. (2024), *Selling Spirals: Avoiding an AI Flash Crash*, Law Fare publication, 8 November
- 9) J.P. Morgan (2023), *How AI will make payments more efficient and reduce fraud*, J.P. Morgan Insights
- 10) Hong Kong Monetary Authority (HKMA) (2025), *HKMA and Cyberport launch second cohort of GenA.I. Sandbox to accelerate A.I. innovation in financial sector*. Press Release, 28 April
- 11) Hu, K. (2023) *ChatGPT sets record for fastest-growing user base – analyst note*, Reuters, 2 February
- 12) IBM (2018), *AI Fairness 360 (AIF360)*
- 13) IndiaAI (2025), *India Takes the Lead: Establishing the IndiaAI Safety Institute for Responsible AI Innovation*, 31 January
- 14) IndiaAI, *Nasscom Responsible AI Resource Kit*
- 15) Indian Banks' Association (IBA) (2025), *IBA Technology Survey and Benchmarking Report*
- 16) Infosys (2025), *The Infosys Responsible AI Toolkit*, February

- 17) International Organization for Standardization and International Electrotechnical Commission (ISO/IEC) (2023) *ISO/IEC 23894:2023 Information Technology – Artificial Intelligence – Guidance on Risk Management*
- 18) ISO/IEC (2023), *ISO/IEC 42001:2023 Information technology – Artificial intelligence – Management system*
- 19) ISO/IEC (2022), *ISO/IEC 23053:2022 Framework for Artificial Intelligence (AI) Systems Using Machine Learning (ML)*.
- 20) Microsoft (2021), *Responsible AI Toolbox*
- 21) McKinsey & Company, (2024), *Scaling gen AI in banking: Choosing the best operating model*
- 22) NITI Aayog (2018), *National strategy for Artificial Intelligence*, NITI Aayog, Government of India
- 23) NITI Aayog (2021), *Responsible AI for All*. NITI Aayog, Government of India
- 24) OECD (2024), *Regulatory approaches to AI in finance*, OECD Artificial Intelligence Papers No. 24, September
- 25) OECD (2019, 2014), *AI Principles*
- 26) Pol, R. and Pachisia, A. (2025), 'Designing India's AI Safety Institute', The Hindu, 5 March
- 27) RBI (2024), *Framework for Responsible and Ethical Enablement of Artificial Intelligence (FREE-AI) in the Financial Sector – Setting up of a committee*, Press Release, 26 December
- 28) SEBI (2025), *Consultation Paper on Guidelines for Responsible Usage of AI/ML in Indian Securities Markets*, 20 June
- 29) Stanford University, (2025), *Artificial Intelligence Index Report*, Stanford Institute for Human Centered Artificial Intelligence.
- 30) Statista (2023), *Global generative AI in finance market size*
- 31) World Economic Forum (WEF) and Accenture (2025), *Artificial Intelligence in Financial Services*. White paper, World Economic Forum, January

Glossary of Key Terms

Sl. No.	Term	Description
1	Accountability	The obligation of individuals or organizations to account for their actions, accept responsibility, and disclose results transparently through specific means and criteria. [ISO/IEC 22989]
2	Adversarial input attacks	Deliberate changes to input data intended to mislead AI models into incorrect decisions or predictions.
3	Agent to Agent (A2A) protocol	A communication protocol enabling autonomous agents to interact without human involvement.
4	Agentic AI	Agentic AI refers to an automated entity that senses and responds to its environment and takes actions to achieve its goals. [ISO/IEC 22989]
5	Model-on-Model risk	Risk arising when one AI system oversees another and fails, potentially causing cascading errors and widespread issues.
6	AI incident	An event where an AI system malfunctions or behaves unpredictably, possibly causing harm or violating safety, fairness or privacy.
7	AI inertia	Resistance within organizations to adopt AI due to cultural, technical, or regulatory barriers.
8	AI inventory	A structured record of all AI systems in use, detailing purpose, risks, dependencies, and performance, to ensure visibility, oversight, and effective risk management across operations.
9	AI Safety Institute	An institution under India AI Mission promoting safe, secure, and trustworthy AI innovation by coordinating research and collaboration across academia, industry, startups, and government.

10	Algorithmic trading	Automated rule-based trading where decisions are made by computer models. [SEBI]
11	Alternative credit scoring models	Using alternative data to assess a borrower's financial health. [HKMA]
12	Artificial Intelligence	An AI system is a machine-based system that, for explicit or implicit objectives, infers, from the input it receives, how to generate outputs such as predictions, content, recommendations, or decisions that can influence physical or virtual environments. Different AI systems vary in their levels of autonomy and adaptiveness after deployment. (Explanatory Memorandum on the updated OECD Definition Of an AI System, December 2023.)
13	Auditability	The ability to inspect and verify system processes and decisions. [maddevs.io]
14	Behaviour Audit	Evaluating AI decisions in real-world settings for ethical and legal alignment.
15	Black box problem	AI systems whose internal workings are opaque to users. [IBM]
16	Credential stuffing	Automated use of stolen credentials to access user accounts. [OWASP]
17	Data minimisation	Collecting only the personal data necessary for a specific service. [ICO]
18	Data poisoning	Manipulating training data to corrupt AI/ML models. [IBM]
19	Decision trees	Hierarchical models used for classification and regression. [IBM]
20	Deepfake	AI-generated or manipulated image, audio or video content that resembles existing persons, objects, places, entities or events and would falsely appear to a person to be authentic or truthful. [EU AI Act, 2024, Article 3(60)]
21	Deep learning	A subset of ML using deep neural networks to mimic human decision-making. [IBM]

22	Differential privacy	A method introducing randomness to protect data privacy without affecting analysis. [IEEE]
23	Dynamic threat detection	Real-time identification and response to cybersecurity threats.
24	Endpoint Detection and Response (EDR)	Software using analytics and automation to protect endpoint devices and IT assets. [IBM]
25	Equity	Fair treatment of individuals. [Merriam-Webster]
26	Estimation and calibration risk	Risk from incorrect or poorly tuned model parameters causing inaccurate outputs.
27	Explainability	Property of an AI system to express important factors influencing the AI system results in a way that humans can understand. [ISO/IEC 22989]
28	Fairness	Ensuring AI decisions are free from undesirable bias or discrimination.
29	Federated learning	Federated learning is a decentralized approach to training machine learning (ML) models. Each node across a distributed network trains a global model using its local data, with a central server aggregating node updates to improve the global model. [IBM]
30	Fine tuning	Adapting pre-trained models for specific tasks. [IBM]
31	Foundation models	Large AI models trained on vast datasets for general tasks. [IBM]
32	Generative AI	Models that generate text, images, or other content. [IBM]
33	GPU (Graphics Processing Unit)	A co-processor designed to accelerate graphics and image processing, and specialized tasks in Machine Learning and Deep Learning involving heavy matrix operations. [IBM]
34	Hallucination	AI hallucination is a phenomenon with Generative AI that produces outputs that are inaccurate and sometimes, non-sensical. [IBM]

35	Homomorphic encryption	Encryption allowing operations on encrypted data by third parties without accessing original data [ISO/IEC 18033-6:2019]
36	Human in the loop/ Human-allied AI	Involving human expertise in the AI lifecycle particularly during training and deployment to actively improve system performance and reliability. [Google Cloud]
37	Inference attacks	Attacks against ML models that infers sensitive attributes of a training data record, given partial knowledge about the record. [NIST]
38	Security information and event management (SIEM)	Application that provides the ability to gather security data and present that data for action via a single interface. [NIST]
39	Landing zones	Scalable cloud configurations for enterprise adoption. [Google Cloud]
40	Large Language Models (LLMs)	Foundation models capable of understanding and generating natural language. [IBM]
41	Logistic regression	A type of linear classifier that predicts the probability of an observation being part of a class. [NIST]
42	Machine learning	A process of optimizing model parameters through computational techniques, such that the model's behaviour reflects the data or experience. [ISO/IEC 22989]
43	Memory scraping	Extracting sensitive data from RAM before encryption, targeting information that is temporarily stored in memory, such as payment card details, login credentials, etc., before it is encrypted or stored persistently.
44	Metadata	Descriptive and structural information about data. (e.g., data format, syntax, and semantics) and descriptive metadata describing data contents (e.g., information security labels). [NIST]

45	Model Bias	<p>Bias- Systematic difference in treatment of certain objects, people or groups in comparison to others [ISO/IEC 22989]</p> <p>Model Bias- Systematic errors in a model arising from erroneous assumptions during the modelling process, that cause it to consistently make incorrect or skewed predictions.</p>
46	Model Concentration	Dependence on a few limited set of models across systems and institutions, increasing systemic risk when these limited models fail.
47	Model Context Protocol (MCP)	An open protocol that standardises how applications provide context to LLMs. [modelcontextprotocol.io]
48	Model Correlation	Statistical similarity between two or more models that measures how similarly or differently the models performs when evaluated on the same datasets.
49	Model Degradation	The gradual or sudden decline in the model's performance over time, resulting in less accurate or reliable predictions compared to its initial performance.
50	Model Distillation	A machine learning technique that aims to transfer the learnings of a large pre-trained model, the "teacher model," to a smaller "student model." [IBM]
51	Model Drift	Model Performance decline due to data or relationship changes. [IBM]
52	Model Herding	Many AI models across institutions behaving similarly due to relying on shared data or design, leading towards model concentration.
53	Model Life Cycle	Life Cycle- evolution of a system, product, service, project or other human-made entity, from conception through retirement. [ISO/IEC 22989]

		Model Life Cycle- The stages an AI model goes through, from creation and training to deployment, monitoring, and eventual retirement
54	Model Manipulation	Intentional alteration or interference with an AI model's parameters, structure, or inputs to influence its behavior or outputs, often for malicious purposes.
55	Model Inversion	Attackers reverse-engineering models to extract data and information [OWASP]
56	Model Risk	The risk of error due to inadequacies in financial risk measurement and valuation models. [ECB]
57	Model Validation	<p>Validation- confirmation, through the provision of objective evidence, that the requirements for a specific intended use or application have been fulfilled [ISO/IEC 22989]</p> <p>Model Validation- Verifying model accuracy and reliability post-training, to ensure the trained AI model performs as per the intended purpose.</p>
58	Natural Language Processing (NLP)	Systems processing and interpreting human language. (language that is or was in active use in a community of people and whose rules are deduced from usage) [ISO/IEC 22989]
59	Neural Networks	A network composed of one or more layers of interconnected simple computing elements, known as neurons, linked by adjustable weights, that processes input data to produce outputs, mimicking the structure and function of neurons in the human brain. [ISO/IEC 22989]
60	Open-source Model	AI models that can be used, examined, altered and distributed, without having to request permission. In some instances, commercial use of the open models may be restricted [IBM]

61	Privileged Access Attacks	An attack that targets privileged accounts such as administrator or root-level accounts to gain unauthorized control of systems, networks, or sensitive data.
62	Prompt Injection	Exploiting prompt concatenation to manipulate AI behaviour [NIST]
63	Red teaming	An exercise, reflecting real-world conditions, that is conducted as a simulated adversarial attempt to compromise organizational missions and/or business processes to provide a comprehensive assessment of the security capability of the information system and organization. [NIST]
64	Retrieval Augmented Generation (RAG)	Combines retrieval systems with Generative AI to deliver contextual responses. RAG systems fetch relevant information from external knowledge bases and provide it to the GenAI model, enabling accurate and up-to-date outputs without retraining the model. [NIST]
65	Small Language Models (SLMs)	AI models, smaller in scope and scale, capable of processing, understanding and generating natural language content, audio, video, etc. [IBM]
66	Support Vector machines	A supervised machine learning algorithm that classifies data by finding an optimal separator that maximizes the distance between each class in an N-dimensional space. [IBM]
67	Synthetic data	Artificial data that is generated from original data and a model that is trained to reproduce the characteristics and structure of the original data. [EDPS]
68	Transparency	Making AI system information made available to relevant stakeholders in a comprehensive, accessible and understandable manner [ISO/IEC 22989:2022(E)]
69	Trinity Models	AI models focused on one language, task, and domain.
70	Understandability	Ease with which users comprehend AI operations and outputs.



FinTech Department, Central Office
Reserve Bank of India
Mumbai